

# System Programming

---

Ion Mandoiu  
Laurent Michel





# Overview

---

- Motivation
- Process Life-cycle
  - Fork
  - Wait
  - Exec
  - File handles
- Pipes
- FIFO



# Motivation

---

- **Processes**

- Historically the core concept in O.S.
- Represent a *running* executable
  - Code + Processor State + Memory + Resource usage


- **Organization in O.S.**

- Processes form a tree rooted at “init”
- Every process has a parent
- Every process has 0 - \* children



```
init—/usr/sbin/apach—8*[ /usr/sbin/apach]
—/usr/sbin/spamd—2*[spamd child]
—acpid
—atd
—bluetoothd
—btsync—6*[{btsync}]
—console-kit-dae—64*[{console-kit-dae}]
—cron
—cupsd
—2*[dbus-daemon]
—dbus-launch
—dhcpcd
—dovecot—
|   —anvil
|   —auth
|   —config
|   —12*[imap]
|   —18*[imap-login]
|   —log
—6*[getty]
—irqbalance
—master—
|   —anvil
|   —pickup
|   —qmgr
|   —tlsmgr
—mdm—mdm—
|   —Xorg—2*[{Xorg}]
|   —mdmwebkit—6*[{mdmwebkit}]
—2*[named—6*[{named}]]
—nmbd
—ntpd
—polkitd—{polkitd}
—rpc.idmapd
—rpc.mountd
—rpc.statd
—rpcbind
—rsyslogd—3*[{rsyslogd}]
—saslauthd—4*[saslauthd]
—smbd—2*[smbd]
—squid3—unlinkd
—ssh—ssh—ssh—ssh—bash—pstree
—udev—2*[udev]
—upstart-file-br
—upstart-socket-
—upstart-udev-br
—xrdp
—xrdp-sesman
```

My shell!





# Process Life-cycle

---

- Birth

- From another process
  - Either a CLI or a GUI
- A process always starts as a *clone* of its parent process
- Then the process *upgrades itself* to running a different executable
- Process *retains access* to the files open by the parent

- Life

- Process can create children processes

- Death

- Eventually calls *exit* or *abort* to “suicide”





# Birth via Cloning

---

- A very simple API to do that

```
#include <unistd.h>

pid_t  fork(void);
```

- Child is an exact copy of the parent
- Semantics
  - In the **parent** process:
    - fork returns the process identifier of the child
    - If a failure occurred, it returns -1 (and sets errno)
  - In the **child** process: fork returns 0 (zero)



# Cloning effect on resources

---

- All the files that were open in the parent....
  - Are accessible and shared in the child!
  - Any operation in parent or child moves the file pointer
- In particular
  - the standard file (in/out/err) are accessible in the child



# Cloning effect on address space

---

- **The parent and the cloned child**
  - Are virtually indistinguishable.
  - All memory is 100% identical.
  - But are distinct copies.
  - Any memory change (stack/heap/static) only affect the caller
  - Thus the parent and his clone can quickly *diverge*





# Concurrency

---

- The parent and the child *both* return from fork
  - This happens **concurrently**
  - Both can run at the same time on a multicore machine
  - You **cannot** assume as to who “returns first”
  - That is true even on a uni-core. [order chosen by OS]



# Usage

---

- Typically

- The parent forks
- When the fork returns, test the return value
  - If zero: We are the child!
  - If  $> 0$ : We are the parent and a child is/will be born!
  - If  $< 0$ : We are the parent and the cloning failed [memory?]
- Branch based on the return value to decide what to do next.



# Example

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t value;
    value = fork();
    printf("In main: value = %d\n",value);
    return 0;
}
```

```
src (master) $ cc fork.c
src (master) $ ./a.out
In main: value = 63689
In main: value = 0
```



# Example

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t value;
    value = fork();
    value = fork();
    printf("In main: value = %d\n",value);
    return 0;
}
```

```
In main: value = 63745
In main: value = 0
In main: value = 63746
In main: value = 0
```



# What should the parent do ?

---

- **Depends on application!**
  - It could run concurrently and ignore the child
  - It could run concurrently and check back on the child later
  - It could wait until the child is done (dies!)



# Example

```
#include <stdio.h>
#include <unistd.h>
#include <time.h>
int fib(int n) {
    if (n<= 1)
        return n;
    else return fib(n-1) + fib(n-2);
}
int main() {
    pid_t value = fork();
    int i;
    if (value == 0) {
        for(i=0;i < 30;i++)
            printf("fib(%2d) = %d\n",i*5,fib(i*5));
    } else {
        long begin = time(NULL);
        for(i=0;i < 10;i++) {
            sleep(1);
            printf("Elapsed time in parent: %ld\n",time(NULL)-begin);
        }
    }
    return 0;
}
```



# What happens ?

---

- **Parent forks**

- On return from fork: test pid value
  - ZERO            This is the child, compute a bunch of fib values.
  - >0            This is the parent, sleep in 1s increment  
Report time at each wakeup  
Then exit

- **When the parent dies...**

- The child lives on until it finishes its loop
- The child has been 'adopted' by an ancestor (typically, init!)



# Waiting on a child

---

- Useful when the child has a task to do
  - Typical of a shell like bash/ksh/zsh/csh/....
- Two simple APIs to do that

```
#include <sys/wait.h>

pid_t    wait(int *stat_loc);
pid_t    waitpid(pid_t pid, int *stat_loc, int options);
```

- Purpose
  - Block the calling process until a child dies
  - Reports in \*stat\_loc
    - the cause of death (check: man wait)
    - the exit status of the child (what he returned from main)
  - return value identifies the dead child






# Please note!

---

- **With Operating Systems**
  - The child usually dies before the parent
  - The parent attends to the dead processes by waiting on them
- **Collecting the dead children...**
  - Is called “reaping” and is the responsibility of the parent process



**Reaper Dreadnought**



# Reaping

---

- When the parent is alive
  - Parent attends to the dead processes by waiting on them
- When the parent is dead
  - Two possibilities
    - Children are adopted by an ancestor process and live on.
      - That's the default.
      - Adopted by “init” (the first process)
    - Children all die in solidarity
      - Precisely: the children are murdered by the dying parent
      - API: `kill(pid_t pid, int sig)`





# How this is done...

---

- Terminating the children ?
  - Loop over child processes and send them kill signals
- Letting init adopt children ?
  - Do nothing, that's the default
- Becoming a (sub)-“reaper” ?
  - A descendant of init can become a Reaper
  - Simply call `prctl(PR_SET_CHILD_SUBREAPER,1);`
  - (man prctl for details)





# Example

```
#include <stdio.h>
#include <unistd.h>
#include <time.h>
#include <sys/wait.h>
int fib(int n) {
    if (n<= 1)
        return n;
    else return fib(n-1) + fib(n-2);
}
int main() {
    pid_t value = fork();
    int i;
    if (value == 0) {
        for(i=0;i < 16;i++)
            printf("fib(%2d) = %d\n",i*3,fib(i*3));
    } else {
        long begin = time(NULL);
        for(i=0;i < 10;i++) {
            sleep(1);
            printf("Elapsed time in parent: %ld\n",time(NULL)-begin);
        }
        int exitStatus;
        pid_t deadChild = wait(&exitStatus);
        printf("Child %d died\n",deadChild);
        printf("Exited normally? [%s] with status %d\n",
            WIFEXITED(exitStatus) ? "yes" : "no",
            WEXITSTATUS(exitStatus));
    }
    return 0;
}
```



# Trace

```
src (master) $ cc -o ffwait fibforkwait.c
src (master) $ ffwait
fib( 0) = 0
fib( 3) = 2
fib( 6) = 8
fib( 9) = 34
fib(12) = 144
fib(15) = 610
fib(18) = 2584
fib(21) = 10946
fib(24) = 46368
fib(27) = 196418
fib(30) = 832040
fib(33) = 3524578
fib(36) = 14930352
fib(39) = 63245986
Elapsed time in parent: 1
Elapsed time in parent: 2
Elapsed time in parent: 3
fib(42) = 267914296
Elapsed time in parent: 4
Elapsed time in parent: 5
Elapsed time in parent: 6
Elapsed time in parent: 7
Elapsed time in parent: 8
Elapsed time in parent: 9
Elapsed time in parent: 10
fib(45) = 1134903170
Child 66772 died
Exited normally? [yes] with status 0
```



# Process Upgrades

---

- Usually....
  - A fresh clone wants to run *different code*
- This can be easily done
  - Child “guts himself”...
    - [meaning it discards the code in its address space]
  - And reloads the process address space with another executable
    - [picked up from the file system of course]
- Note
  - Opened files are *NOT AFFECTED* by the operation.



# The `exec` family

---

- The act of ‘gutting and upgrading’ is done by the child with a system call
  - There is a whole family (variants) of system calls

```
#include <unistd.h>

extern char **environ;

int execl(const char *path,
          const char *arg0, ...
          /*, (char *)0 */);
```

- Check
  - `man -S3 execl` for variants



# Fundamentally

---

- **The call takes as input**
  - The path to the executable to load inside our own address space
  - A list of arguments to be passed to the new executable
  - A final NULL pointer to give the “end of argument list”
- **Note**
  - If successful **exec1** does not return!
  - Instead, control is transferred to the main function of the new executable!





# Exec example

First a child executable

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char* argv[]) {
    int i, value, sum=0;
    for(i=1; i<argc; i++)
        sum += atoi(argv[i]);
    printf("sum is: %d\n", sum);
    return 0;
}
```

This is a simple “adder” program that computes the sum of its integer arguments



# Parent Program

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <string.h>
#include <errno.h>
int main() {
    pid_t child = fork();
    if (child == 0) {
        printf("In child!\n");
        execl("./adder", "./adder", "1", "2", "3", "10", NULL);
        printf("Oops.... something went really wrong. Shouldn't be here!\n");
        char* msg = strerror(errno);
        printf("Error was: %s\n", msg);
        return -1;
    } else {
        printf("In parent!\n");
        execl("./adder", "./adder", "100", "200", "300", NULL);
        printf("Oops.... something went really wrong. Shouldn't be here!\n");
        char* msg = strerror(errno);
        printf("Error was: %s\n", msg);
        return -1;
    }
}
```



# Behavior

---

- Quite straightforward

```
src (master) $ ./paradder  
In parent!  
In child!  
sum is: 600  
sum is: 16
```

- Imagine typo in name of executable

```
In parent!  
Oops.... something went really wrong. Shouldn't be here!  
Error was: No such file or directory  
In child!  
sum is: 16
```



---

**FORK / EXEC**

**AND**

**INPUT/OUTPUT IMPACTS**

(or what happens to file descriptors)

A SysProg Saga



# Inheritance

---

- Whenever a parent *forks*....
  - The child inherits all the parent's file descriptor....
  - Illustration!



# Playing with redirections

---

- **Note**

- A child process inherits any files opened by the parent

- **Corollary**

- This can be handy to change where the child...
  - Reads its input from
  - Write its output to!

- **How?**

- Simply *change* the files corresponding to
  - stdin (0) | stdout (1) | stderr (2)



# Shell redirections

---

- That's the embryo of piping

- Available in your shell

```
$ sort < file.txt > sorted.txt
```

- Simple syntax

- **<** filename : Take input from the file *filename*
- **>** filename : Send output to the file *filename*
- **2>** filename : Send errors to the file *filename*
- Implemented with the close/open/dup technique



Demo  
Time!





# Brief aside

---

- Remember the IO APIs
  - The “f” family (fopen, fclose, fread, fgetc, fscanf, fprintf,...)
  - All these use a FILE\* abstraction to represent a file
  - All these rely on buffering in the C library
  - Support line-ending translation when opening in “text mode”
- UNIX has a lower-level API for file handling
  - Directly mapped to system calls
  - No buffering
  - Raw IO
  - Uses OS level *file descriptors* **[these are just integers]**



# Brief aside

- APIs highlight (only some of them)

```
#include <fcntl.h>
#include <unistd.h>

int open(const char *path, int oflag, ...);
int close(int fildes);

ssize_t read(int fildes, void *buf, size_t nbyte);
ssize_t write(int fildes, const void *buf, size_t nbyte);

int dup(int fildes);
int dup2(int fildes, int fildes2);
```

“New” APIs



# Brief aside


---

- The API is perfectly fine for
  - Binary file
  - And setting up redirections ;-)
- Two new “special” APIs
  - dup
  - dup2
- Purpose
  - They duplicate file handles



# Brief aside (picture!)

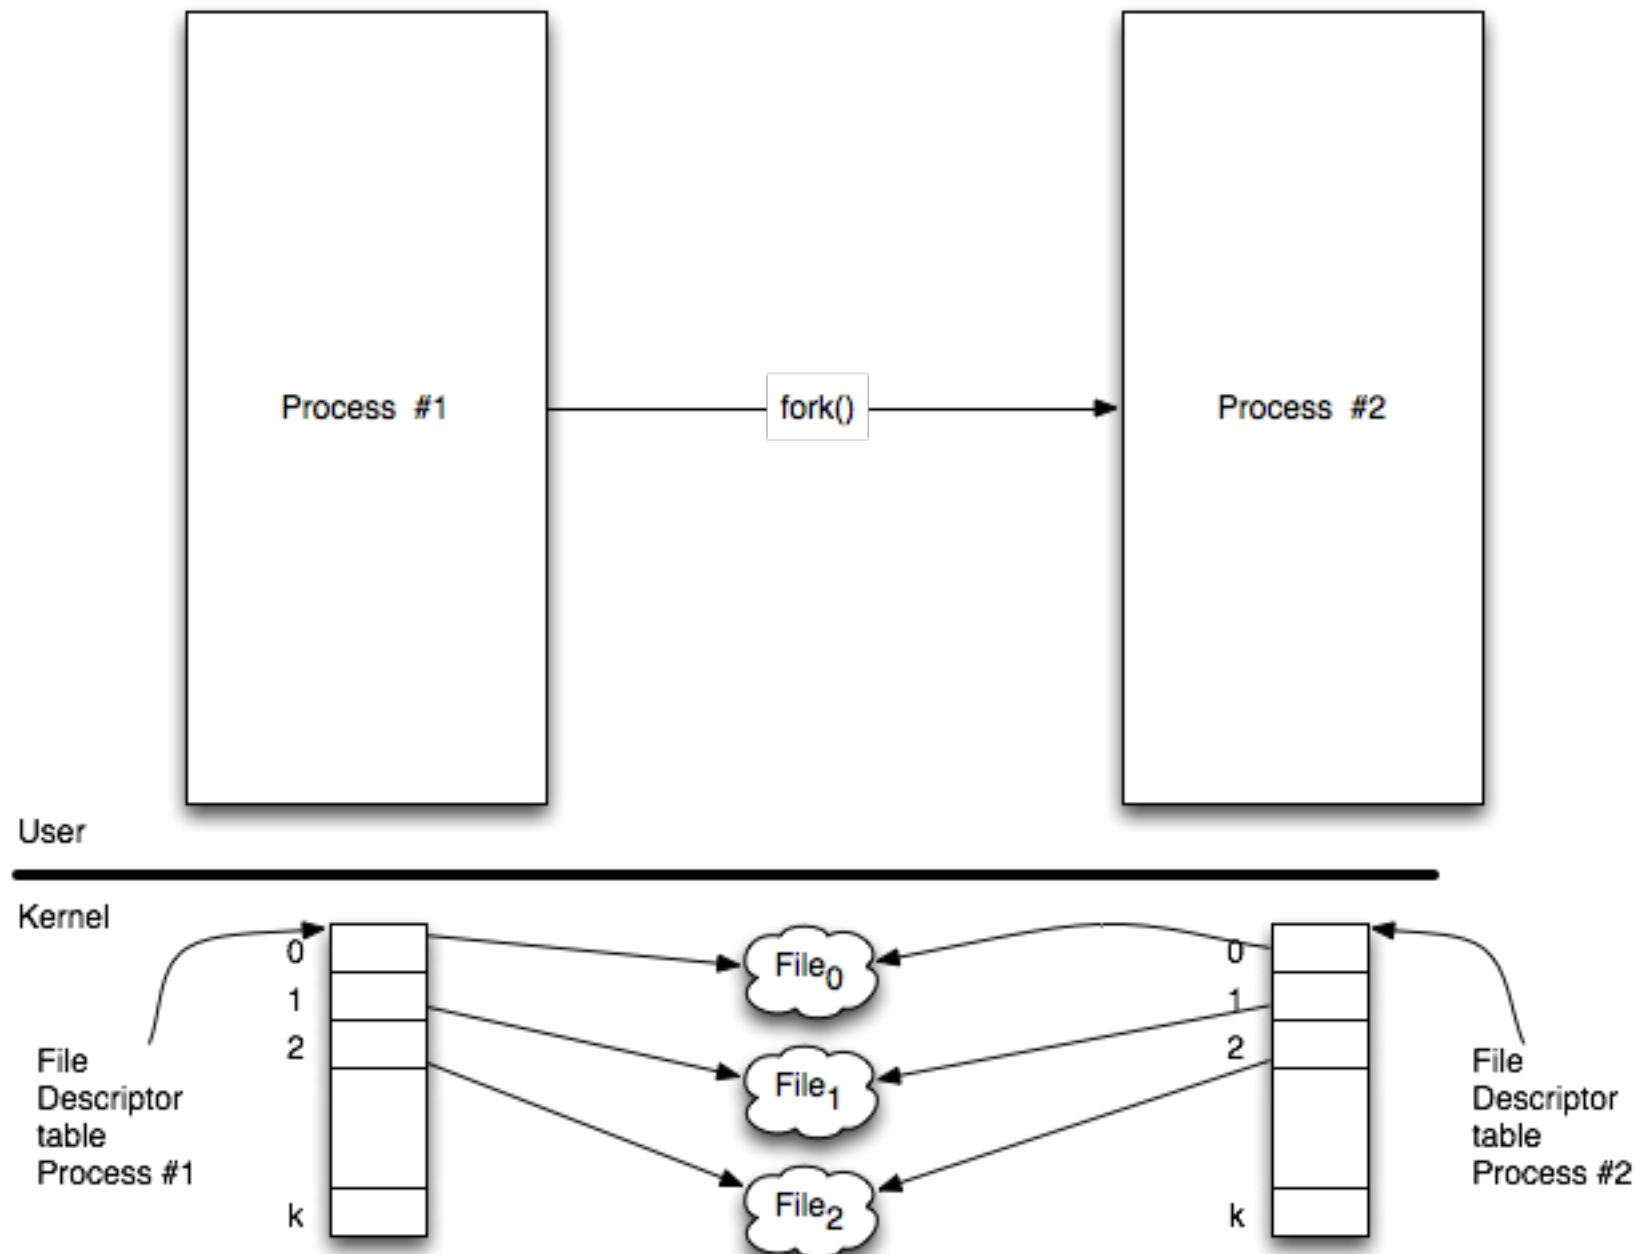
## Two processes in general





# Brief aside (picture!)

Two processes, process #2 is a fork of #1

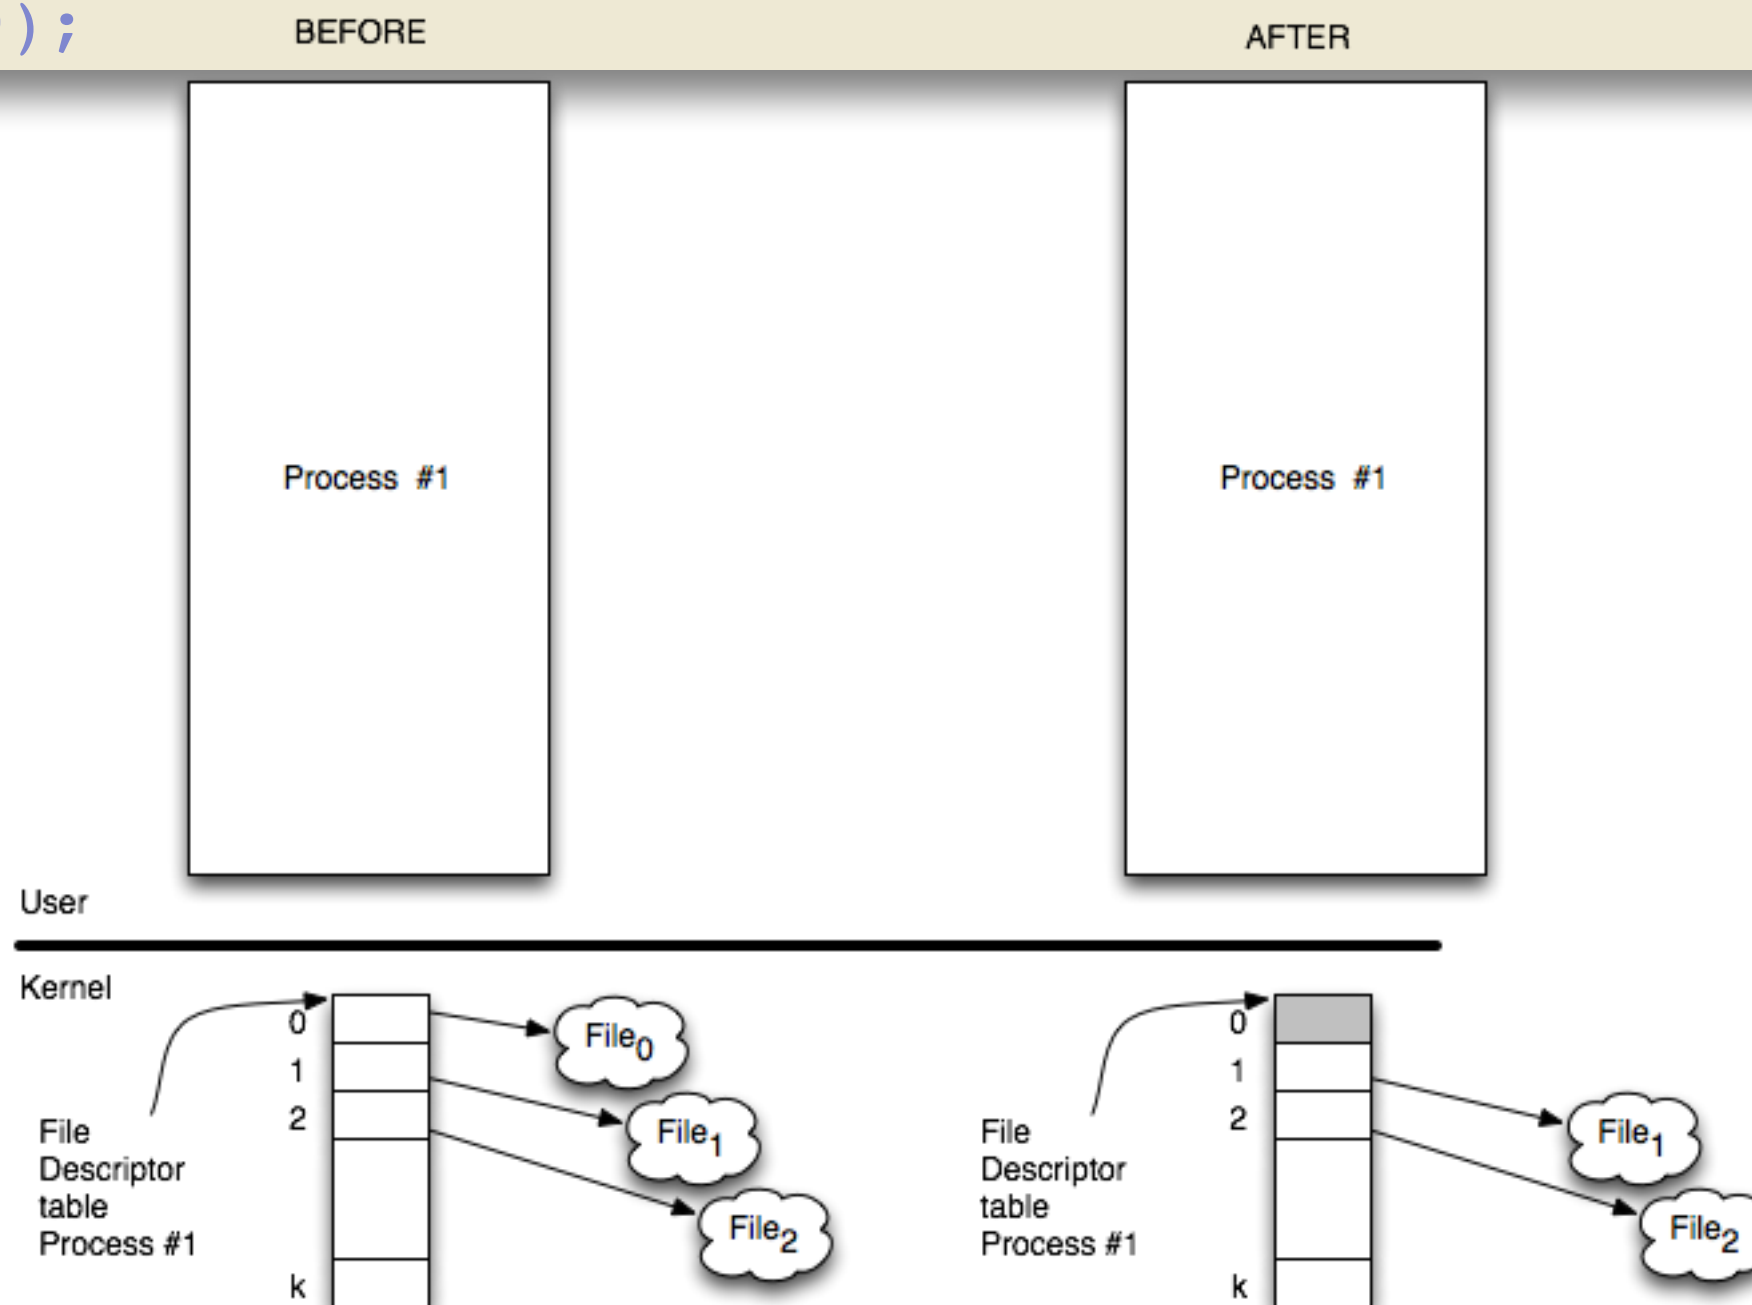




# What close does

- Consider

```
close(0);
```

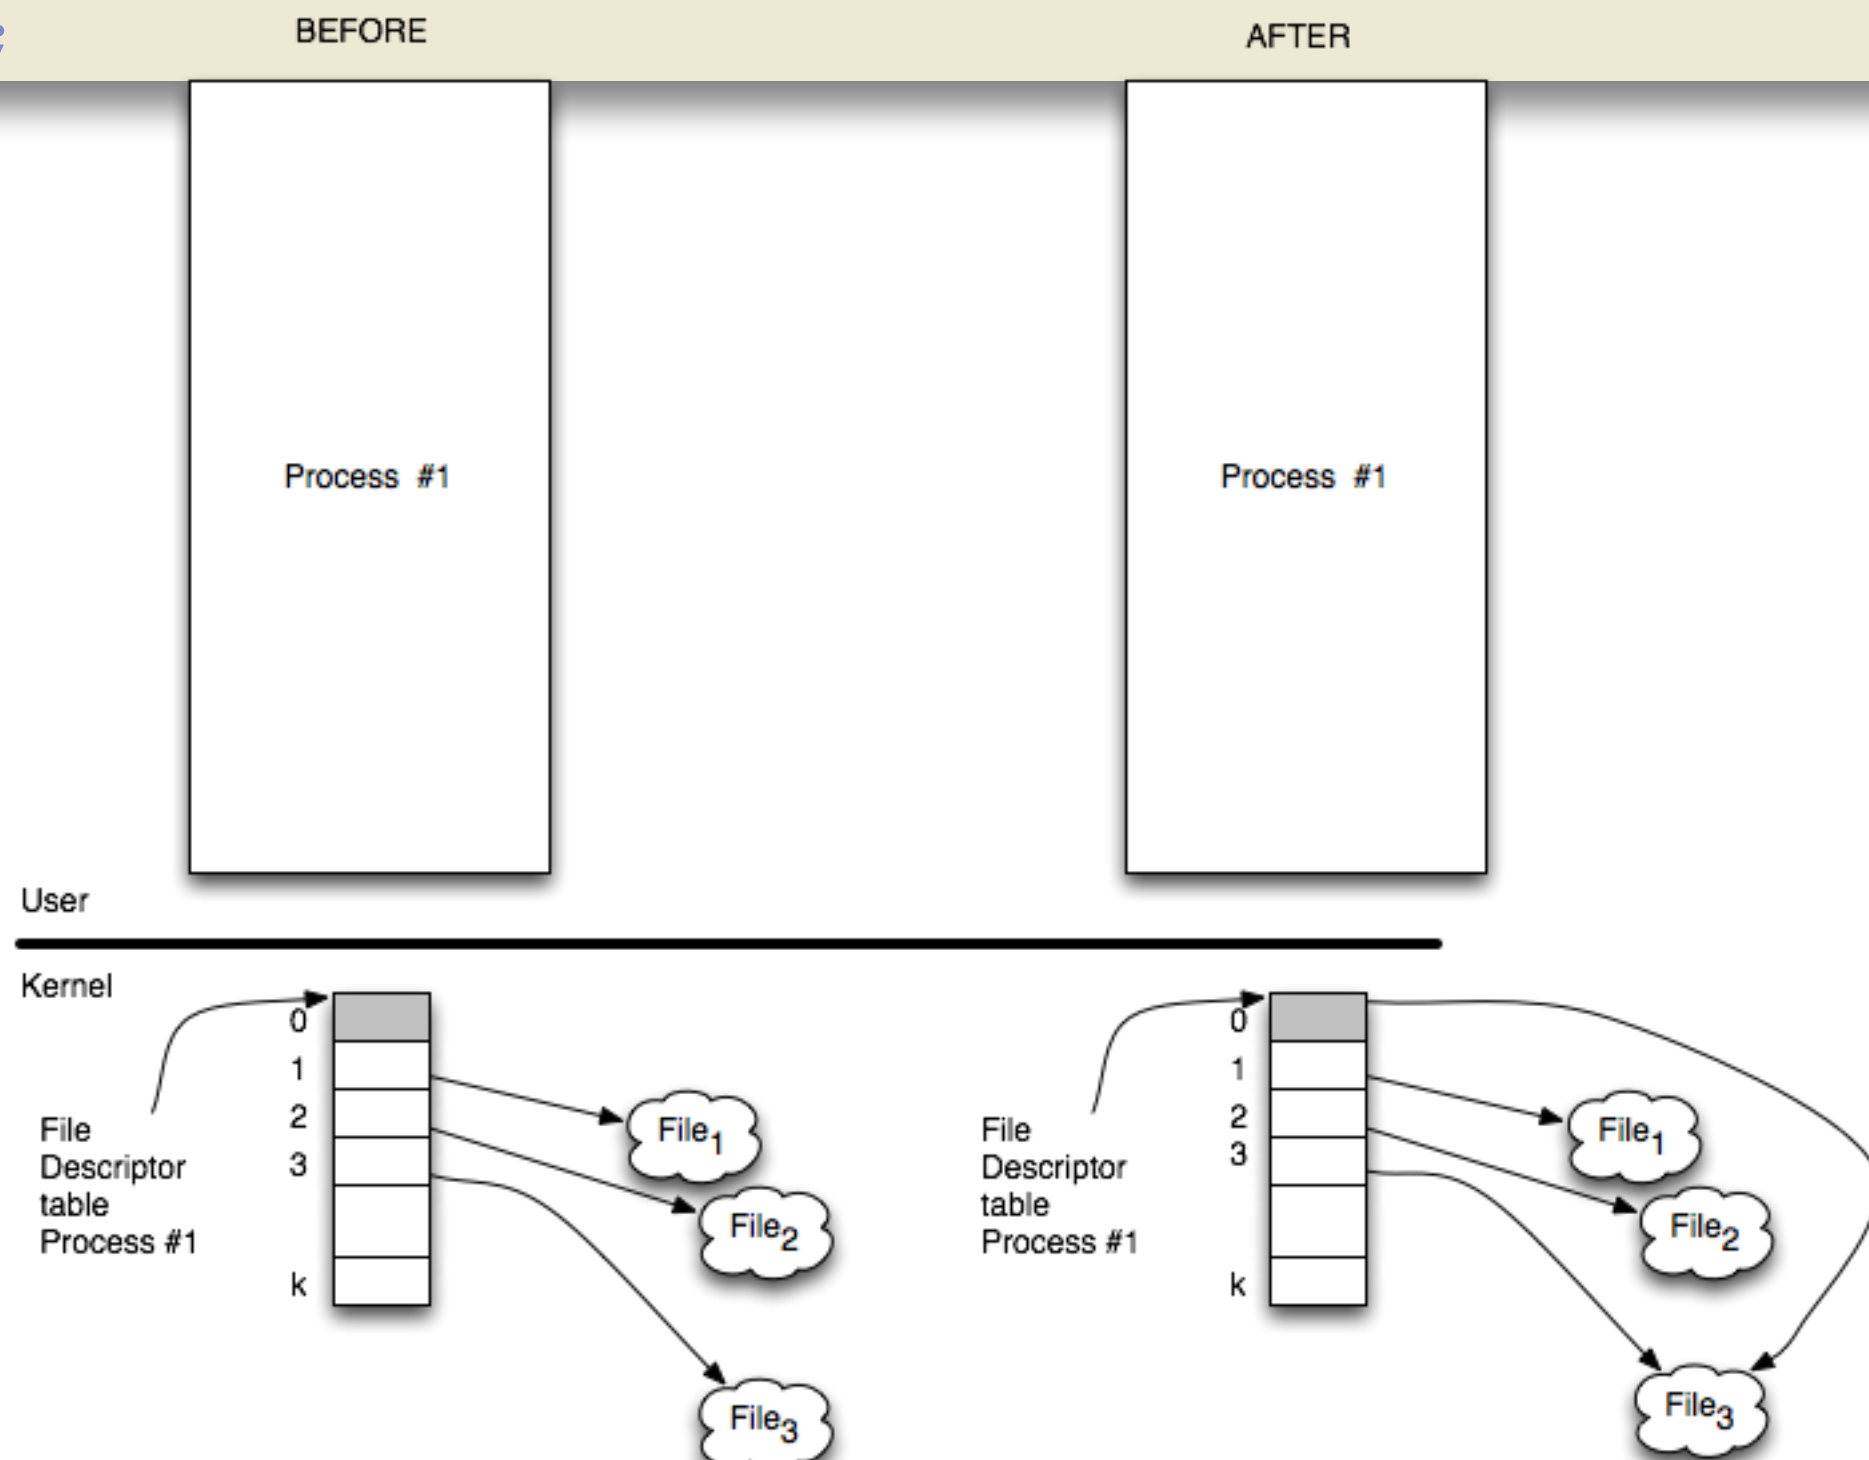




# What dup does

- Duplicate a file handle!

```
dup( 3 );
```





# Bottom line

---

- Purpose of dup

- Ensure that specific entries of the FDT point to the “correct” file

- Remember

- STDIN     0
- STDOUT   1
- STDERR   2





# Key questions

---

- Who should call dup?
  - The parent?
  - The child?
- When should dup be called ?
  - before execl ?
  - after execl ?



# What if?

---

- You call `close(0)` before forking....
  - Then the child will inherit a closed file as STDIN
  - But the parent *also* loses its STDIN!
- You call `close(0)` after forking....
  - Then the child inherited the parent's STDIN
  - The parent's STDIN is unaffected
  - The clone can close its own and reopen the right file before `exec`
- Can you call `close(0)` after `exec` ?
  - No! That's too late!

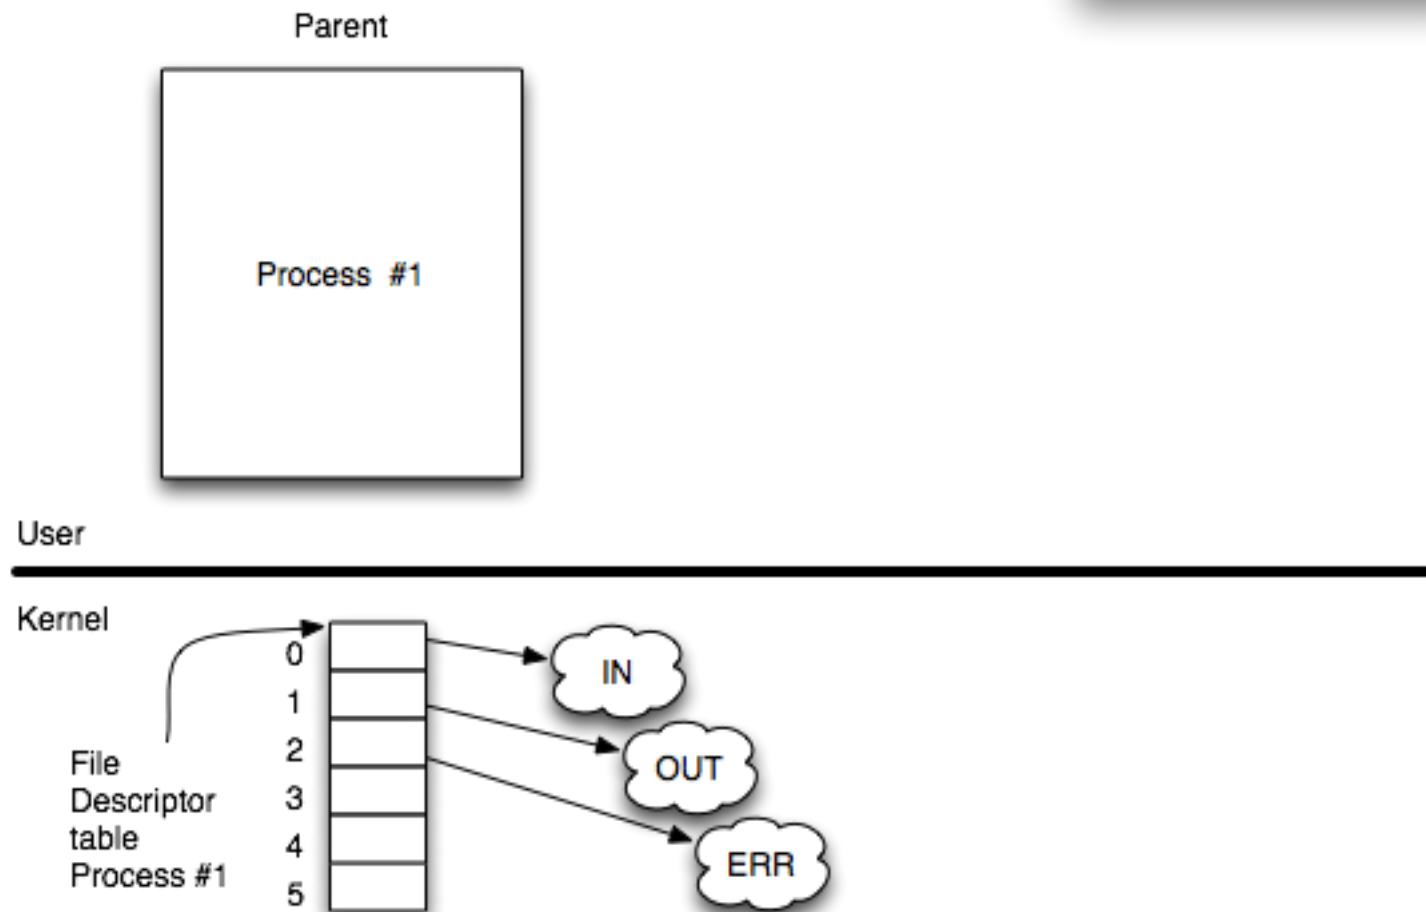
Pipe Demo

Time!



# Pipes

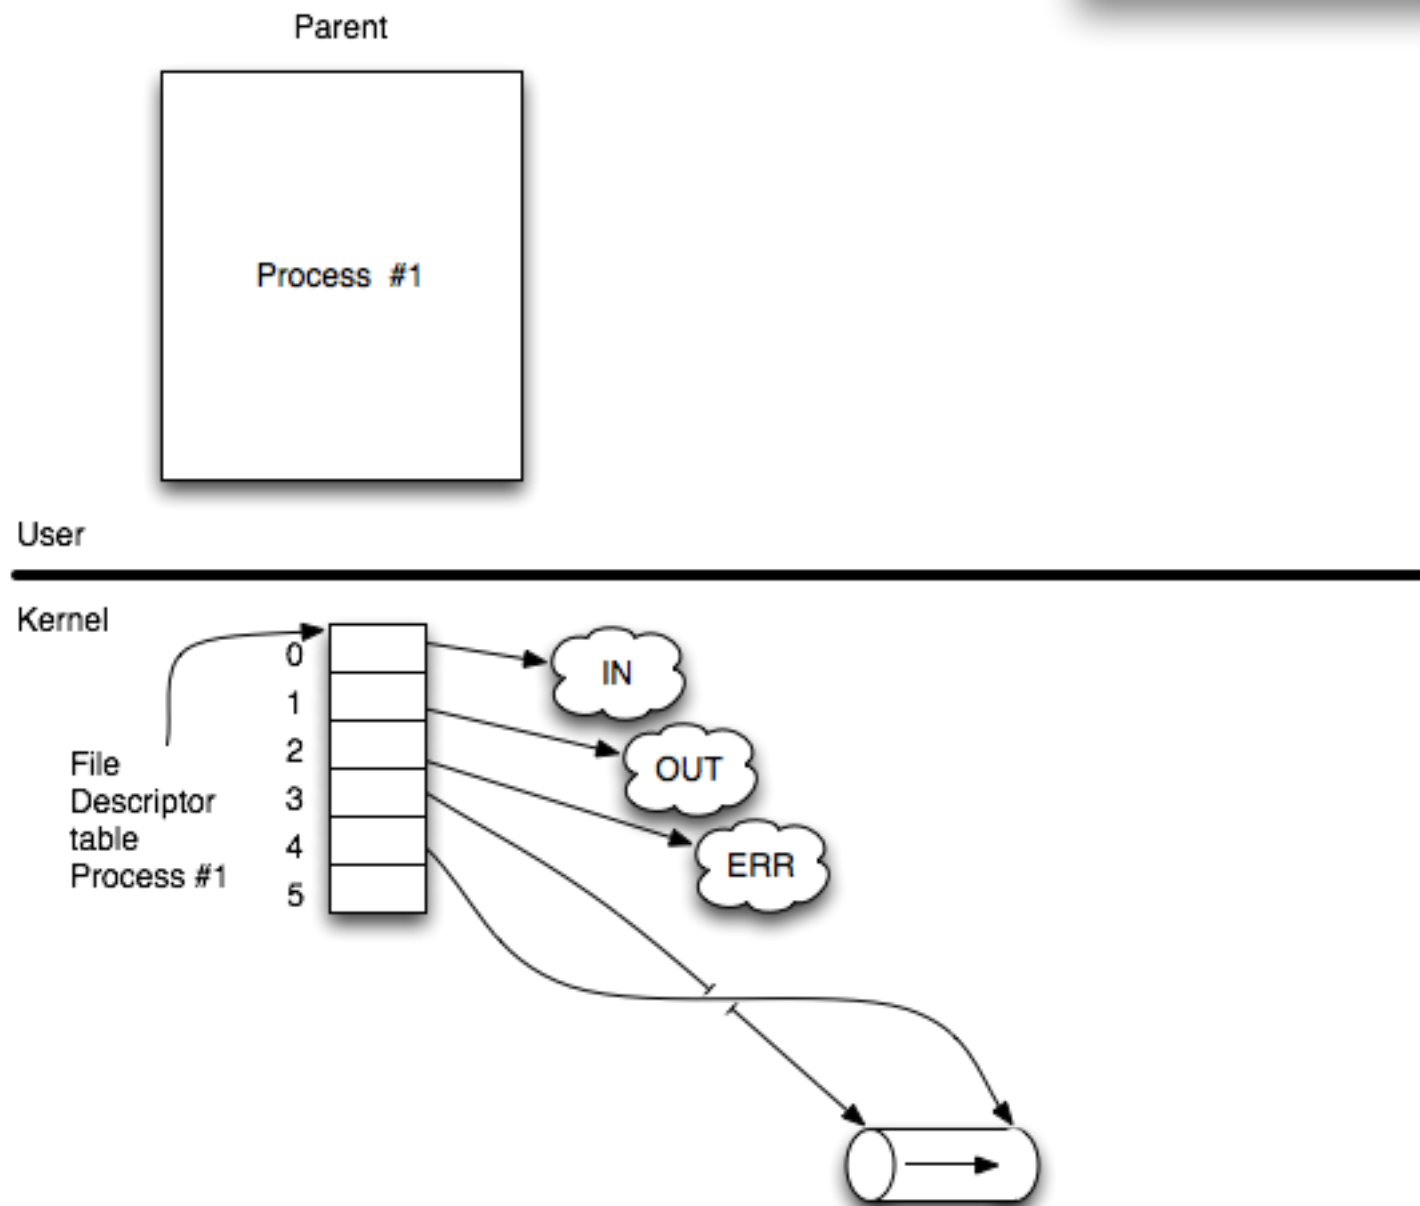
\$ A | B





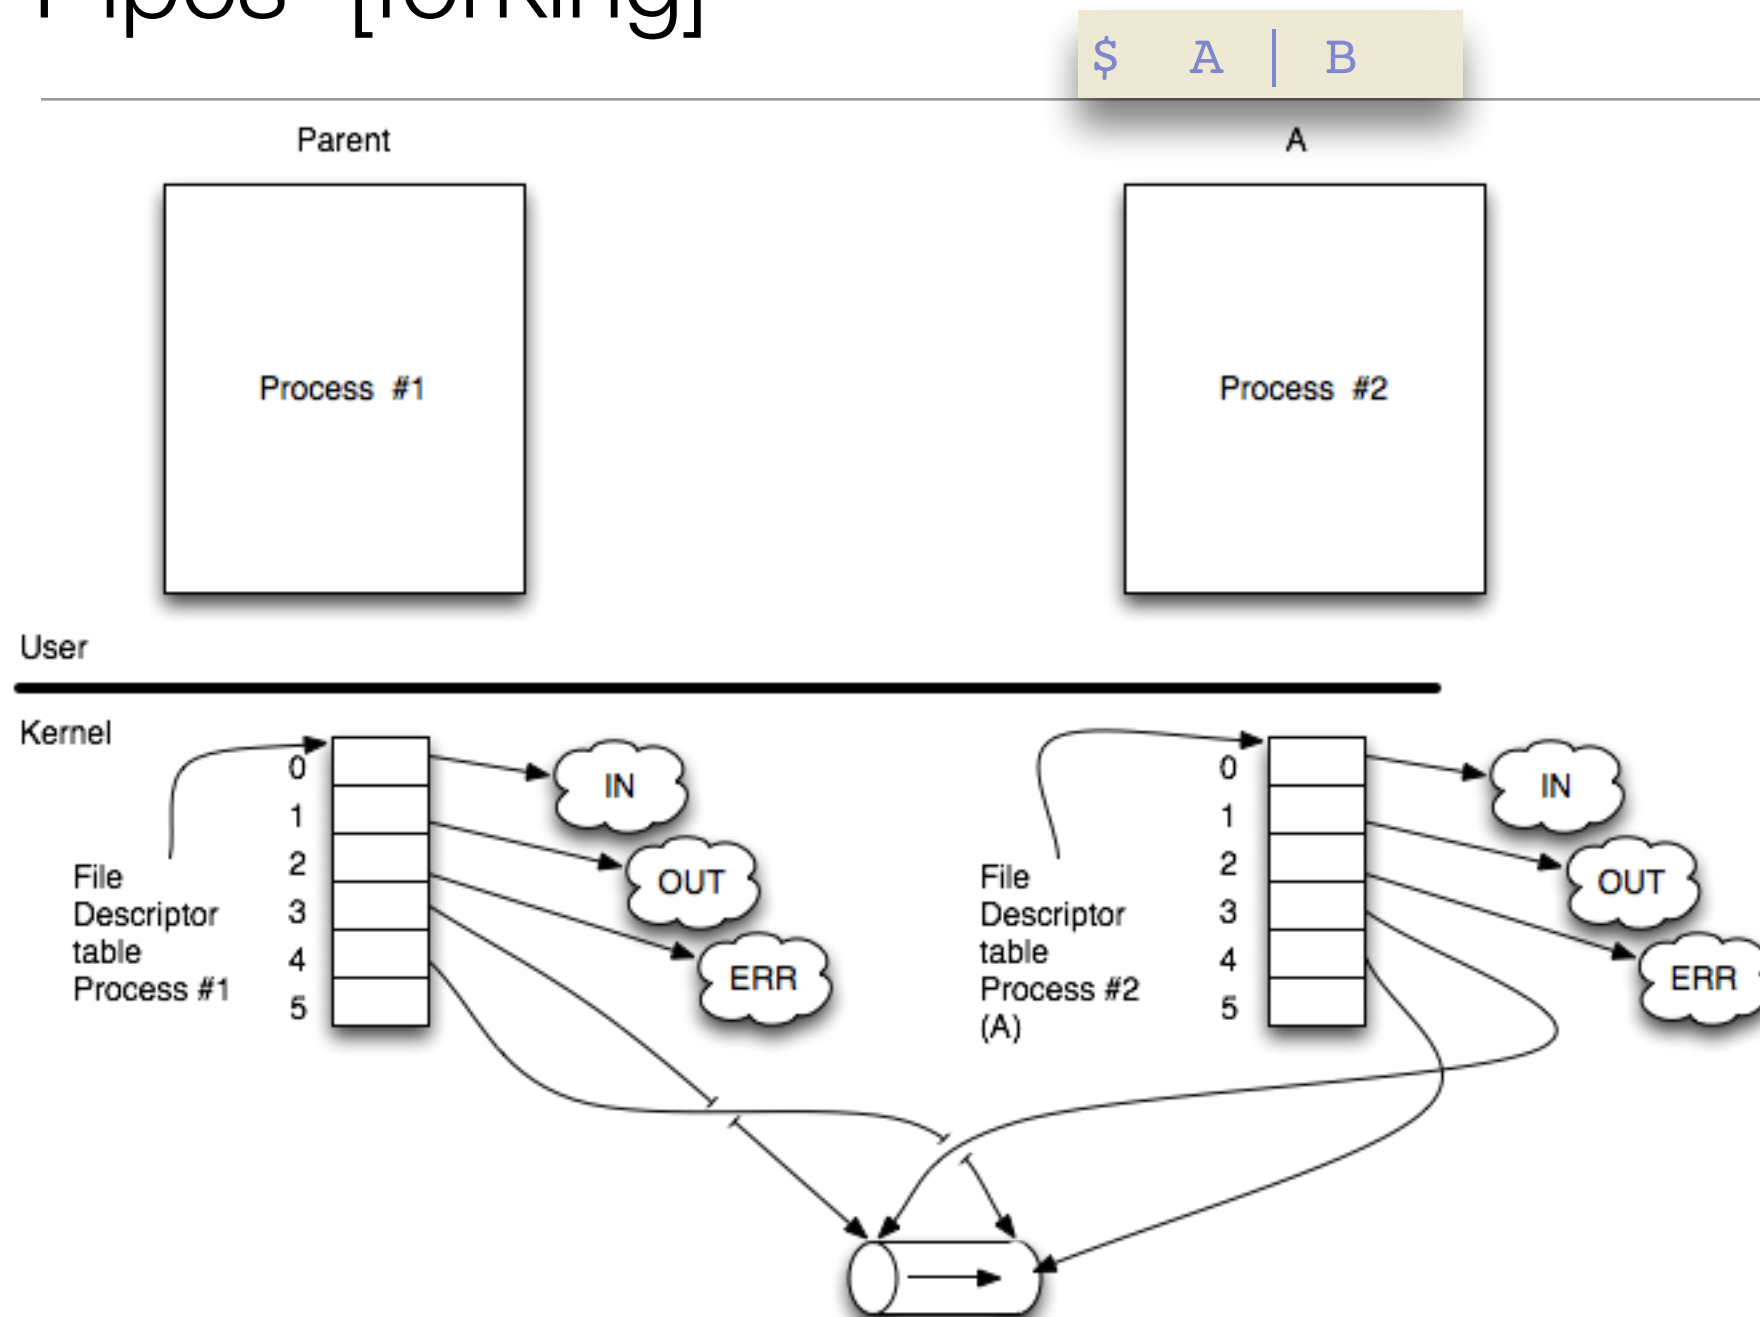
# Pipes [creation]

\$ A | B





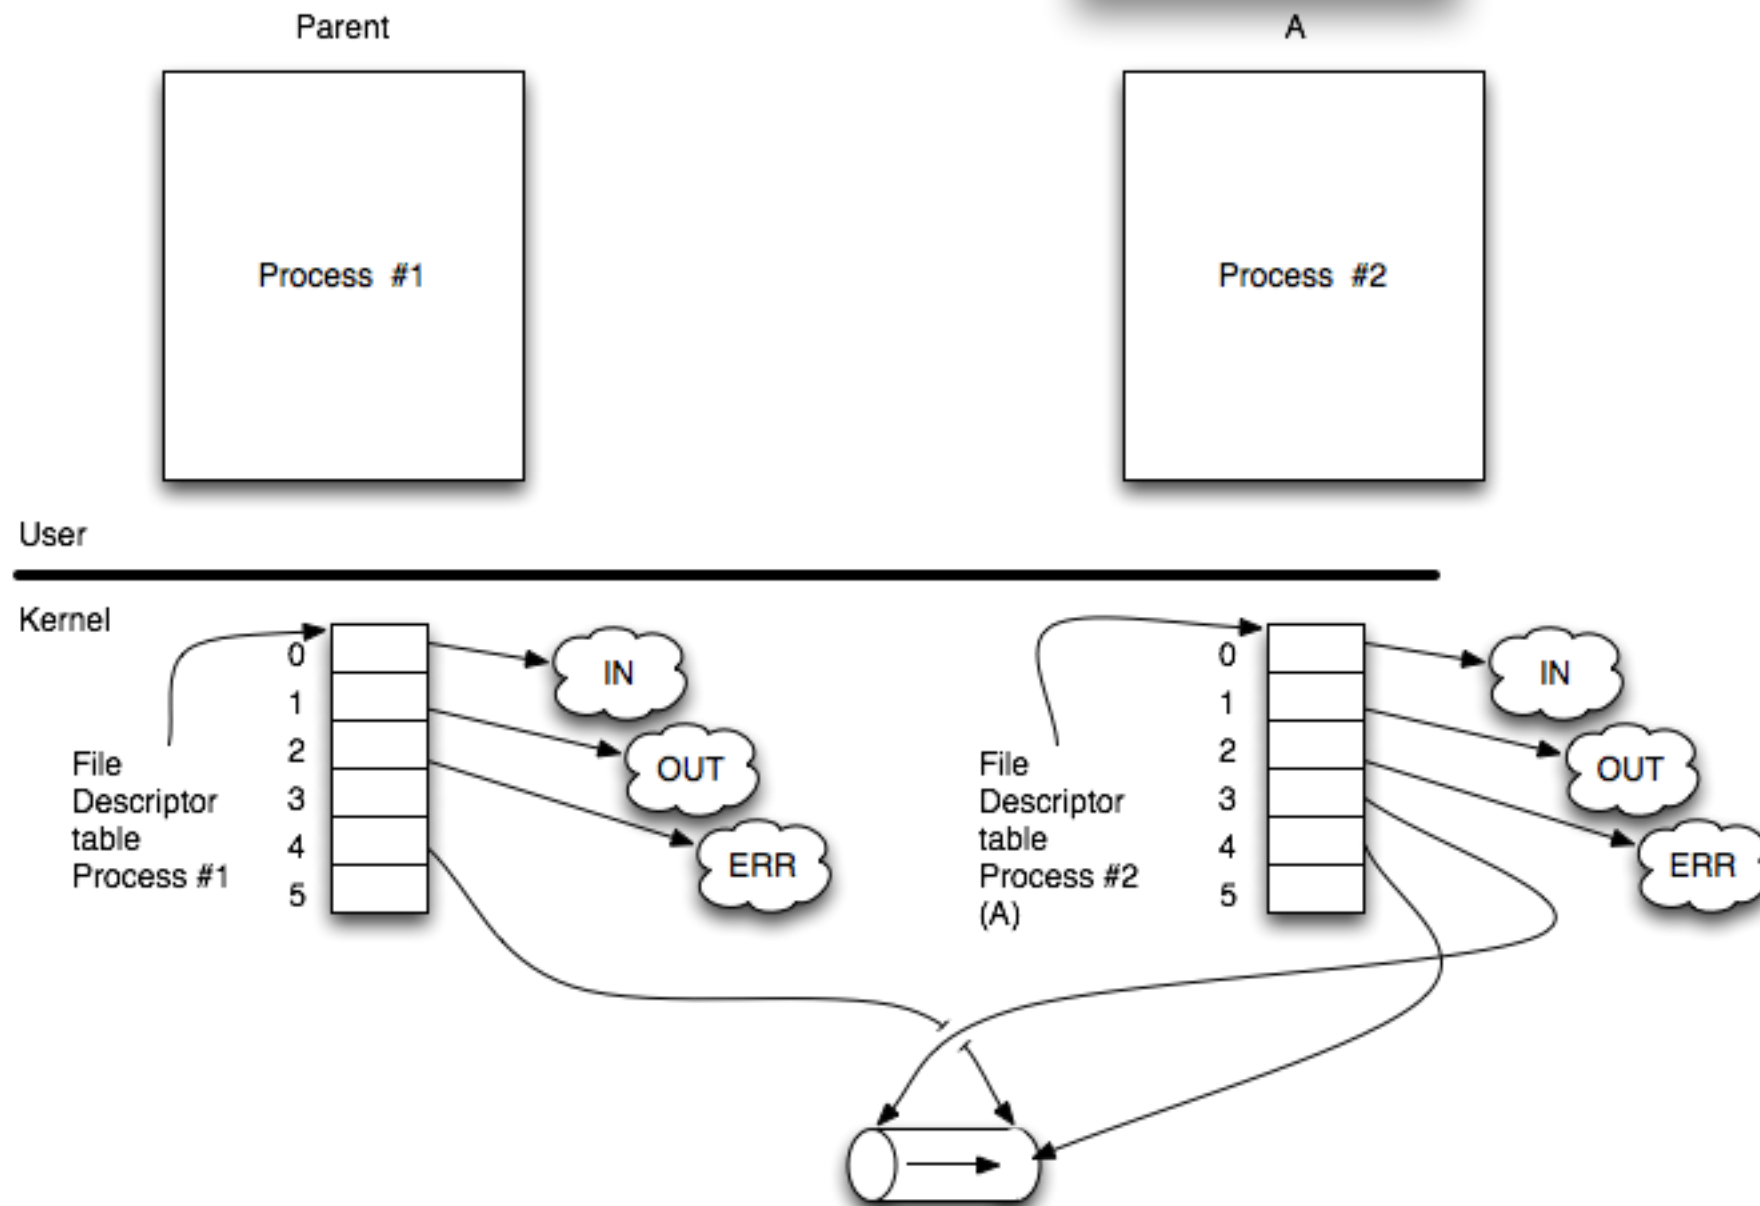
# Pipes [forking]





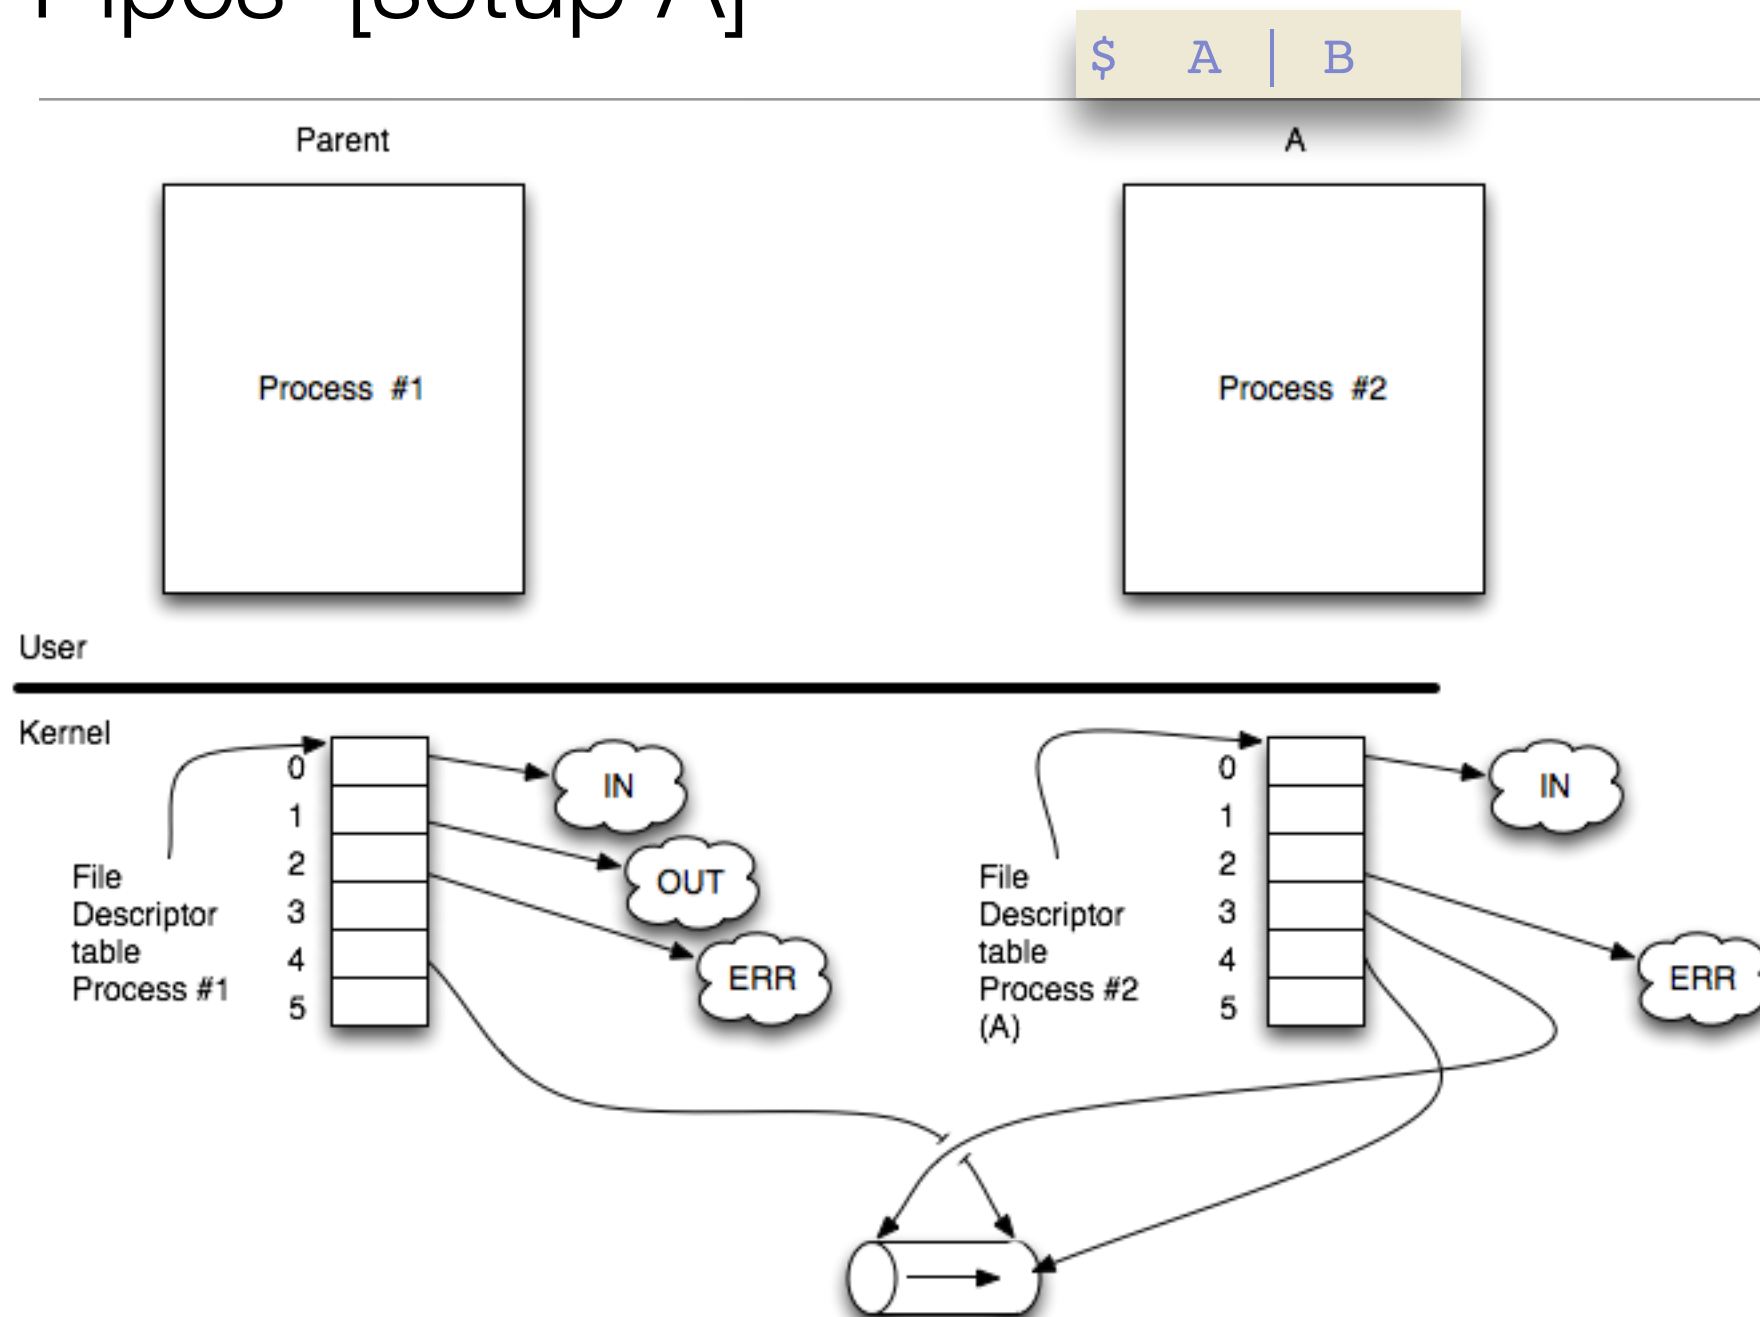
# Pipes [cleaning part 1]

\$ A | B





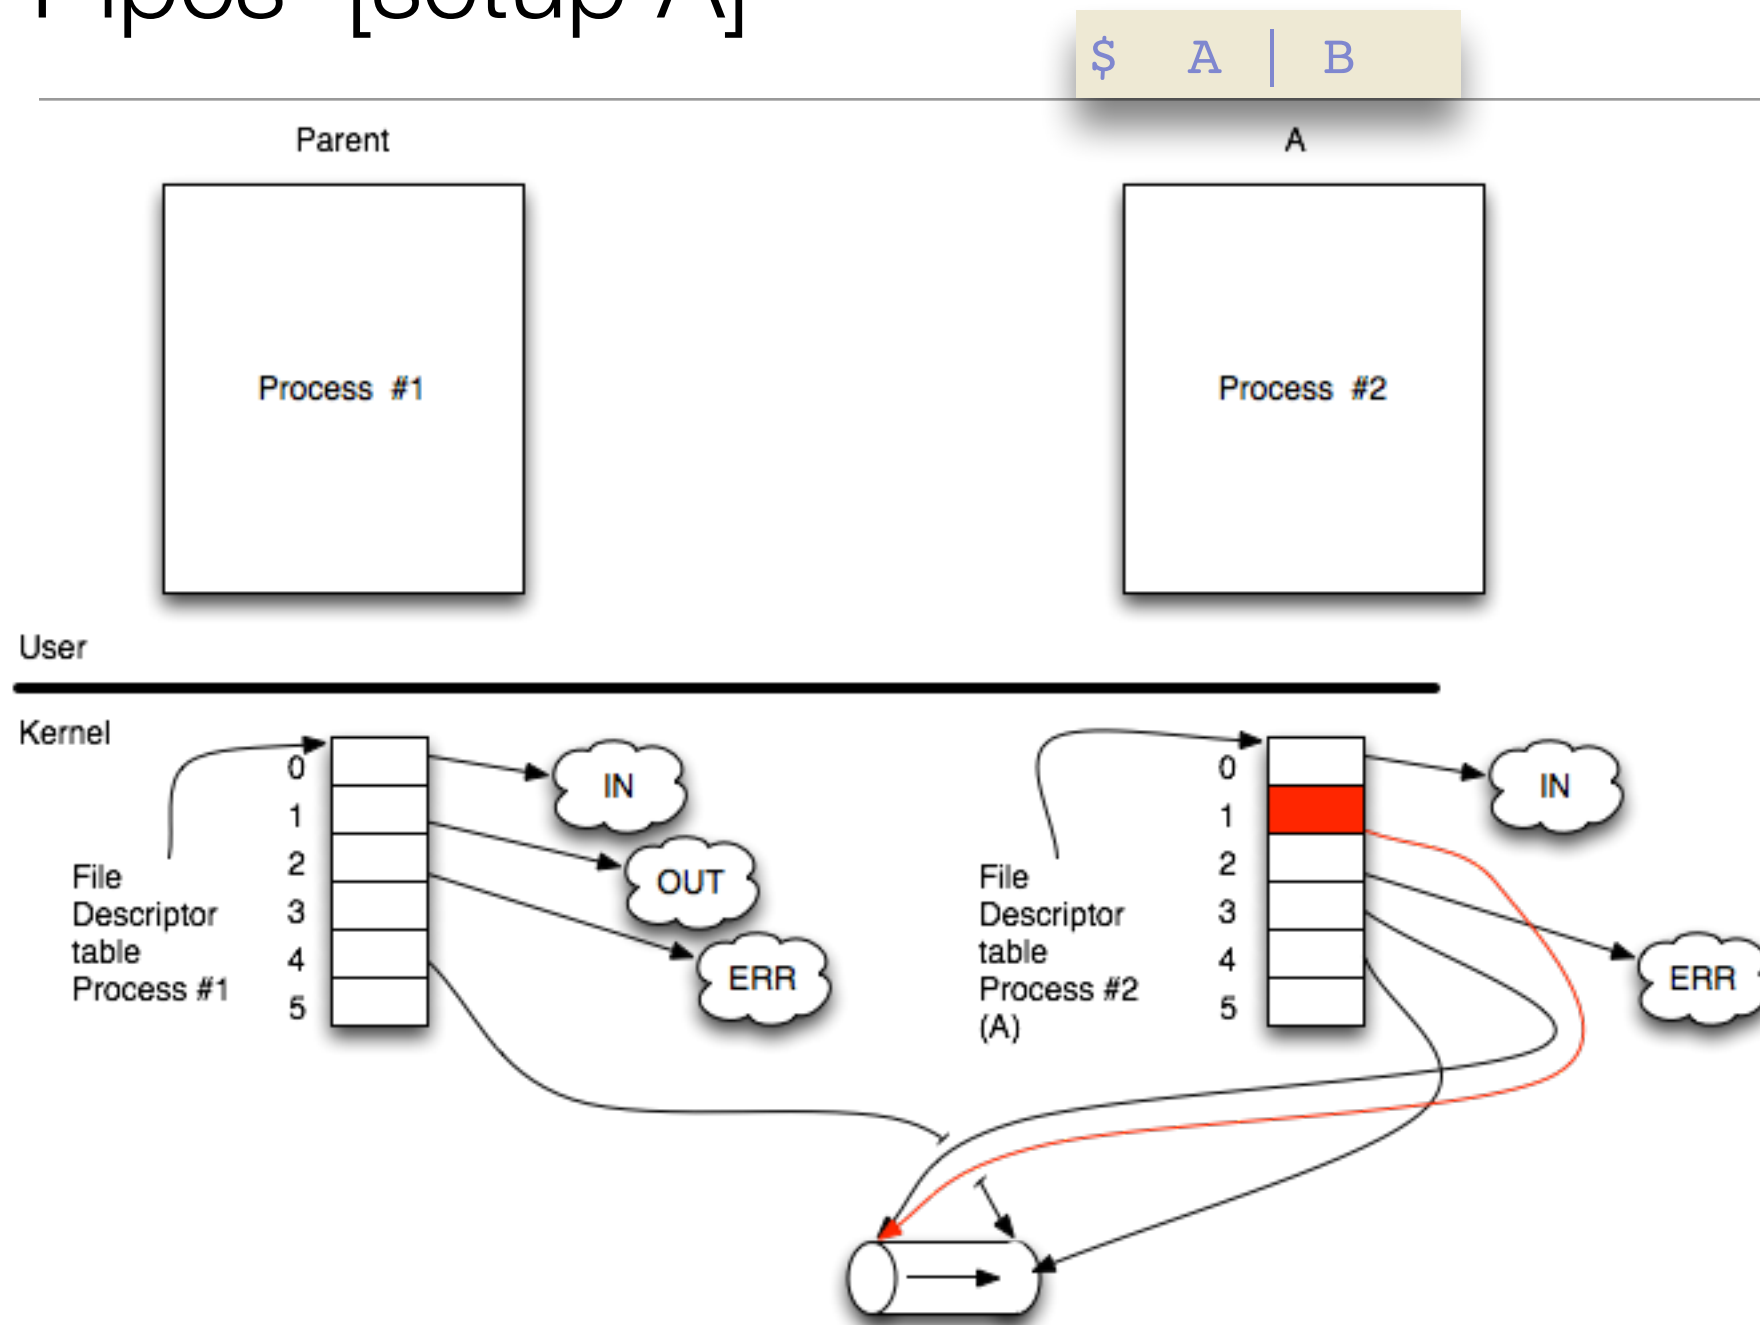
# Pipes [setup A]





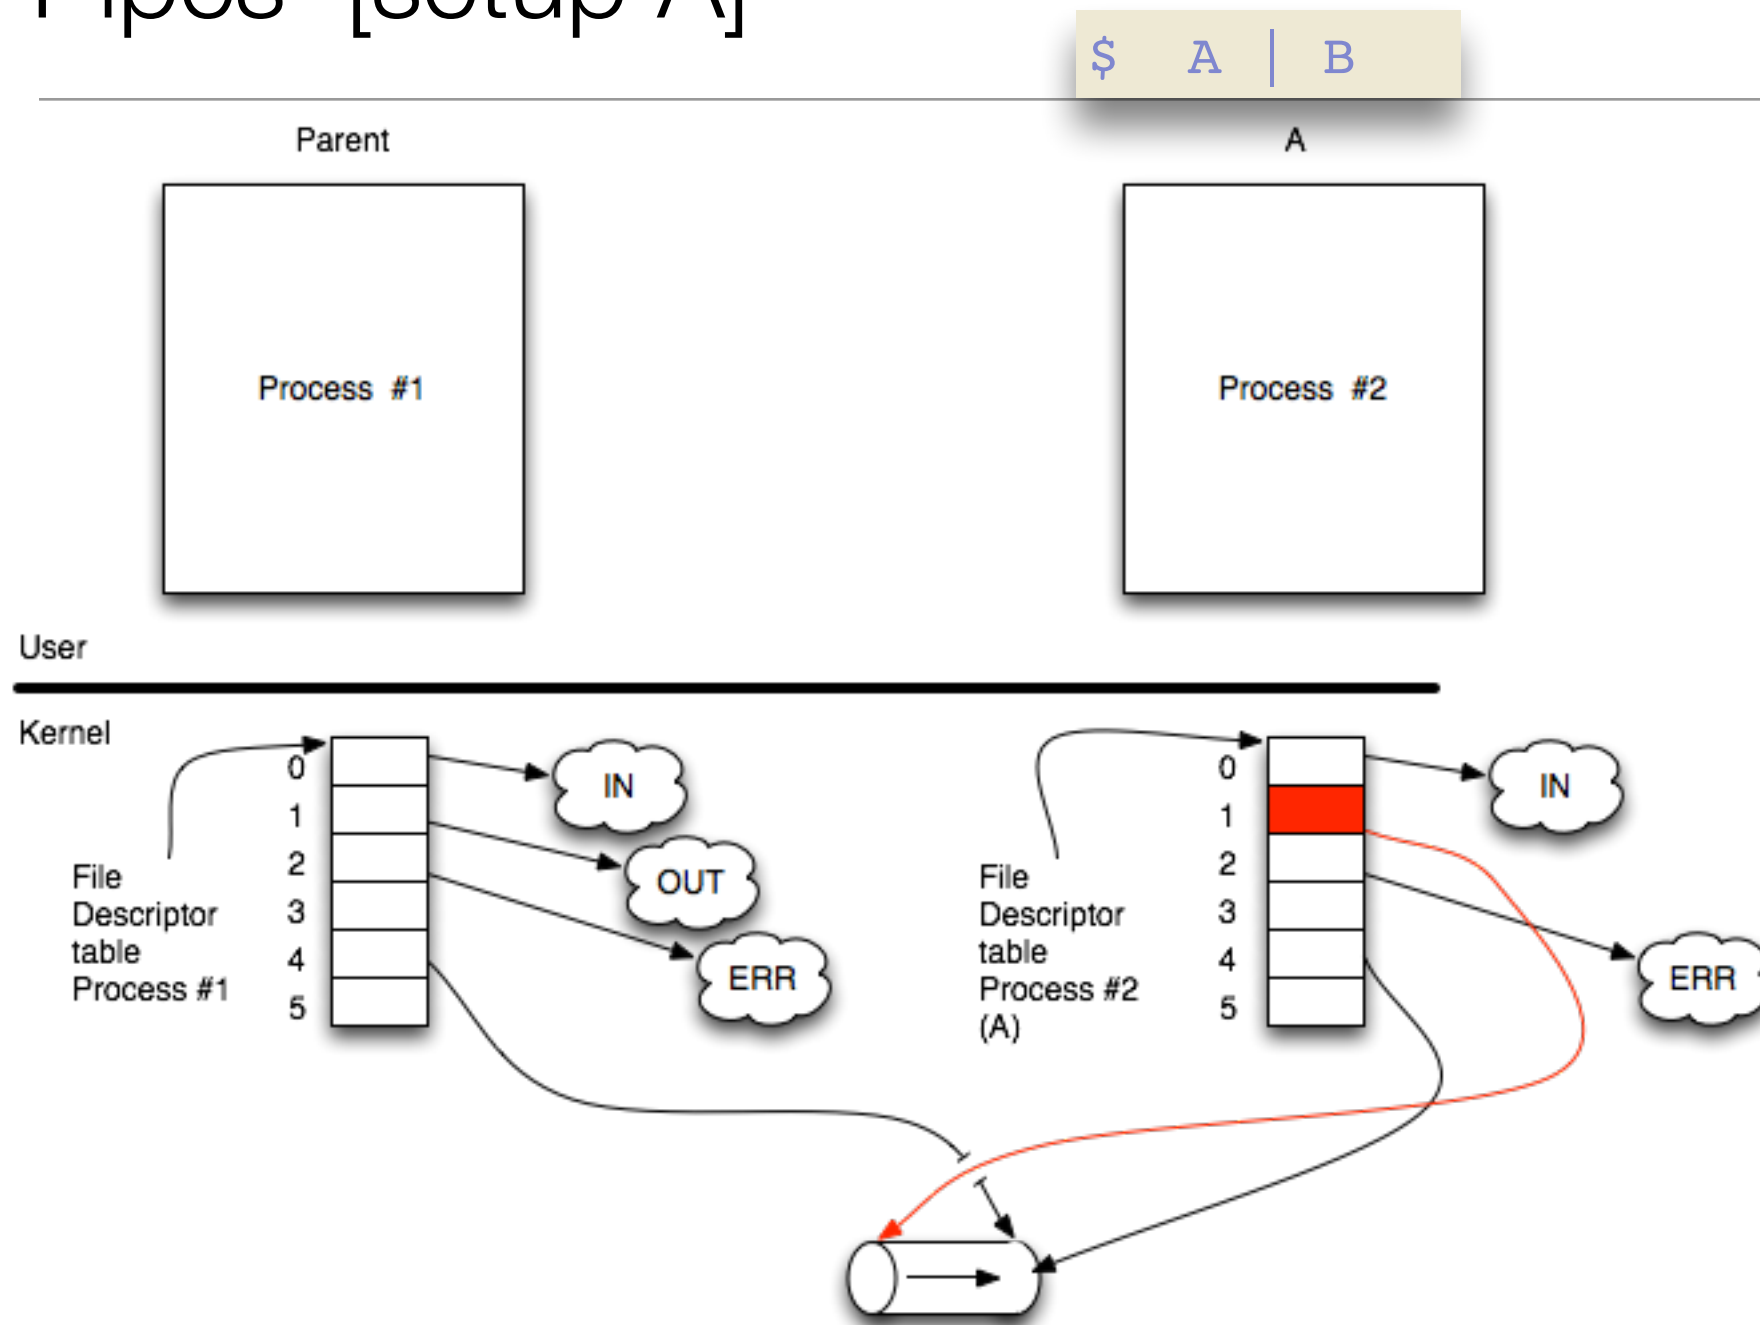


# Pipes [setup A]



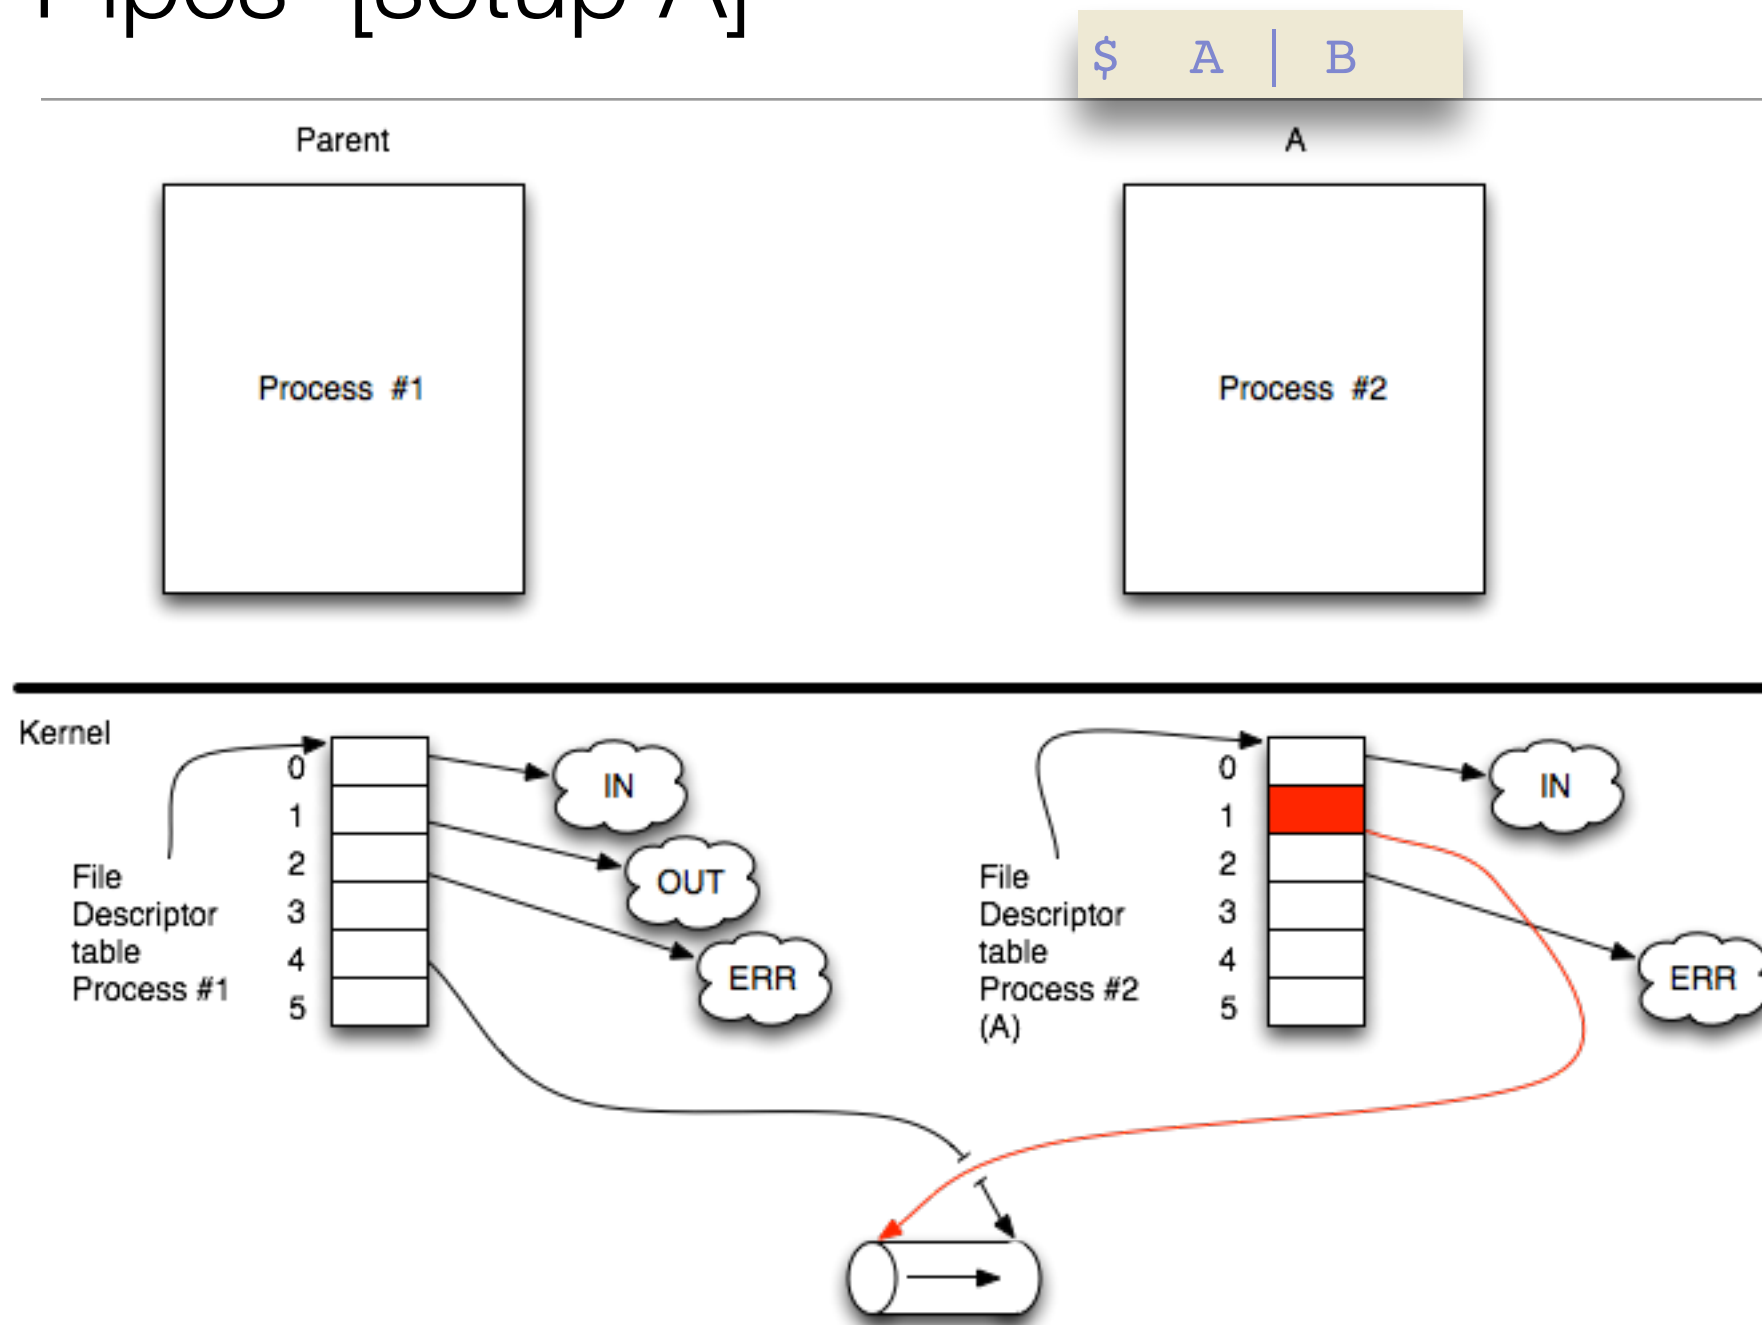


# Pipes [setup A]





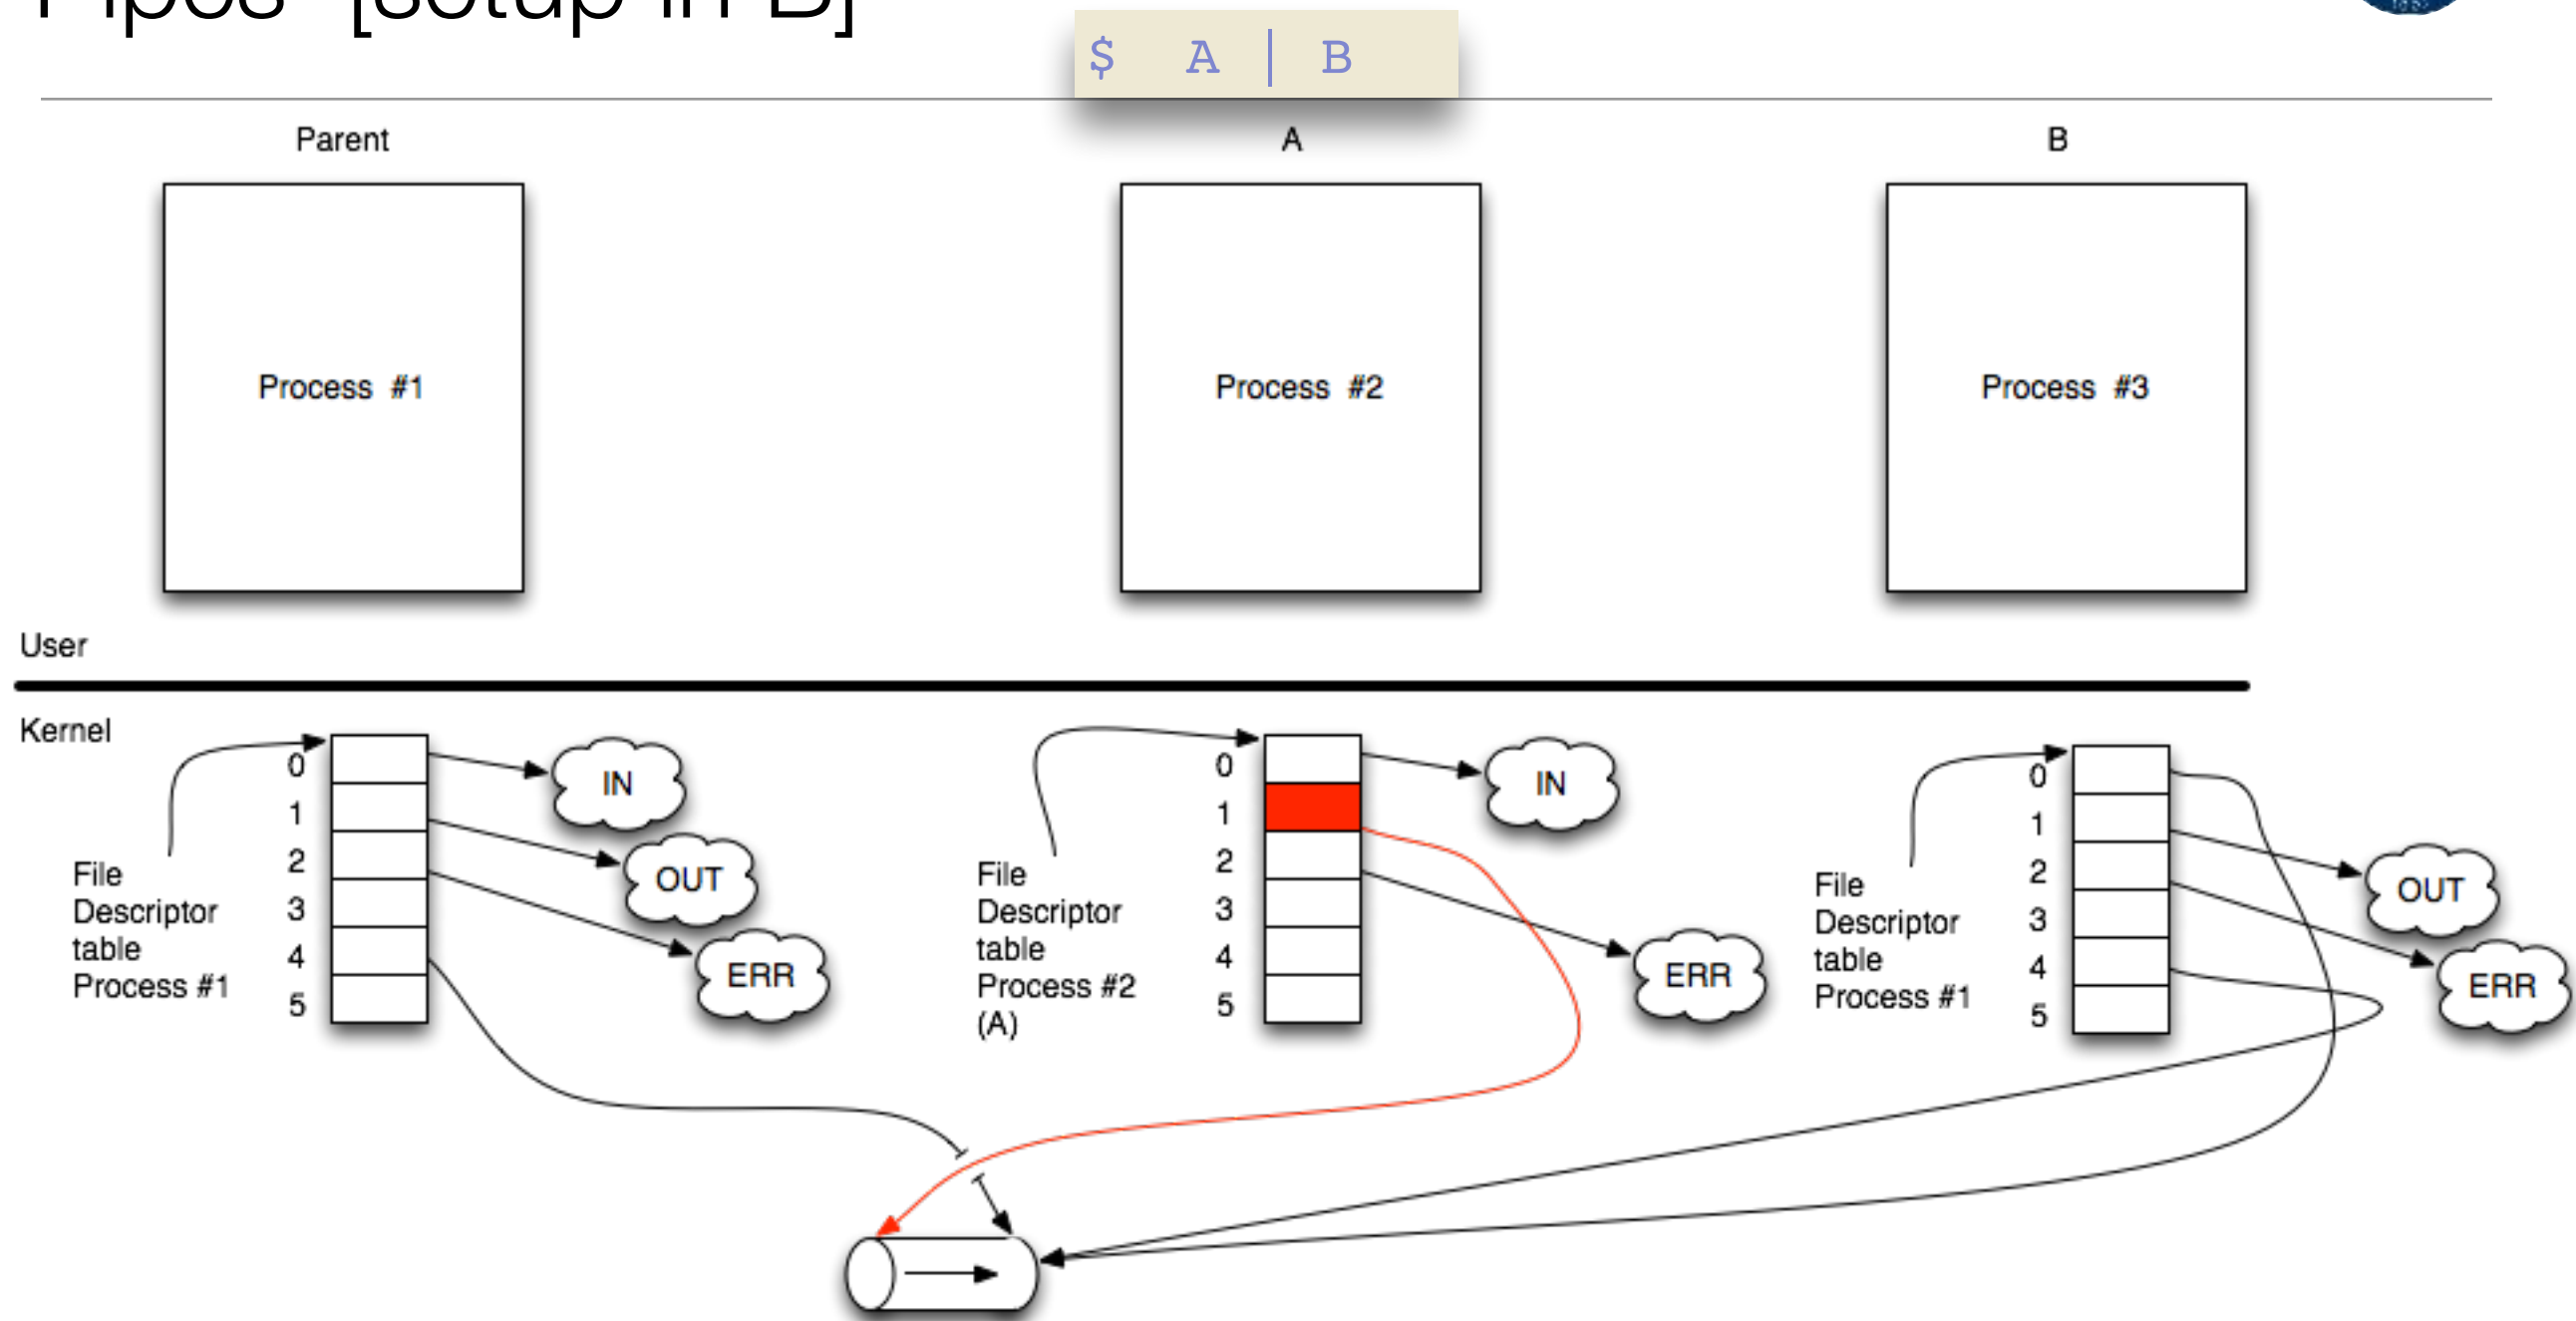
# Pipes [setup A]





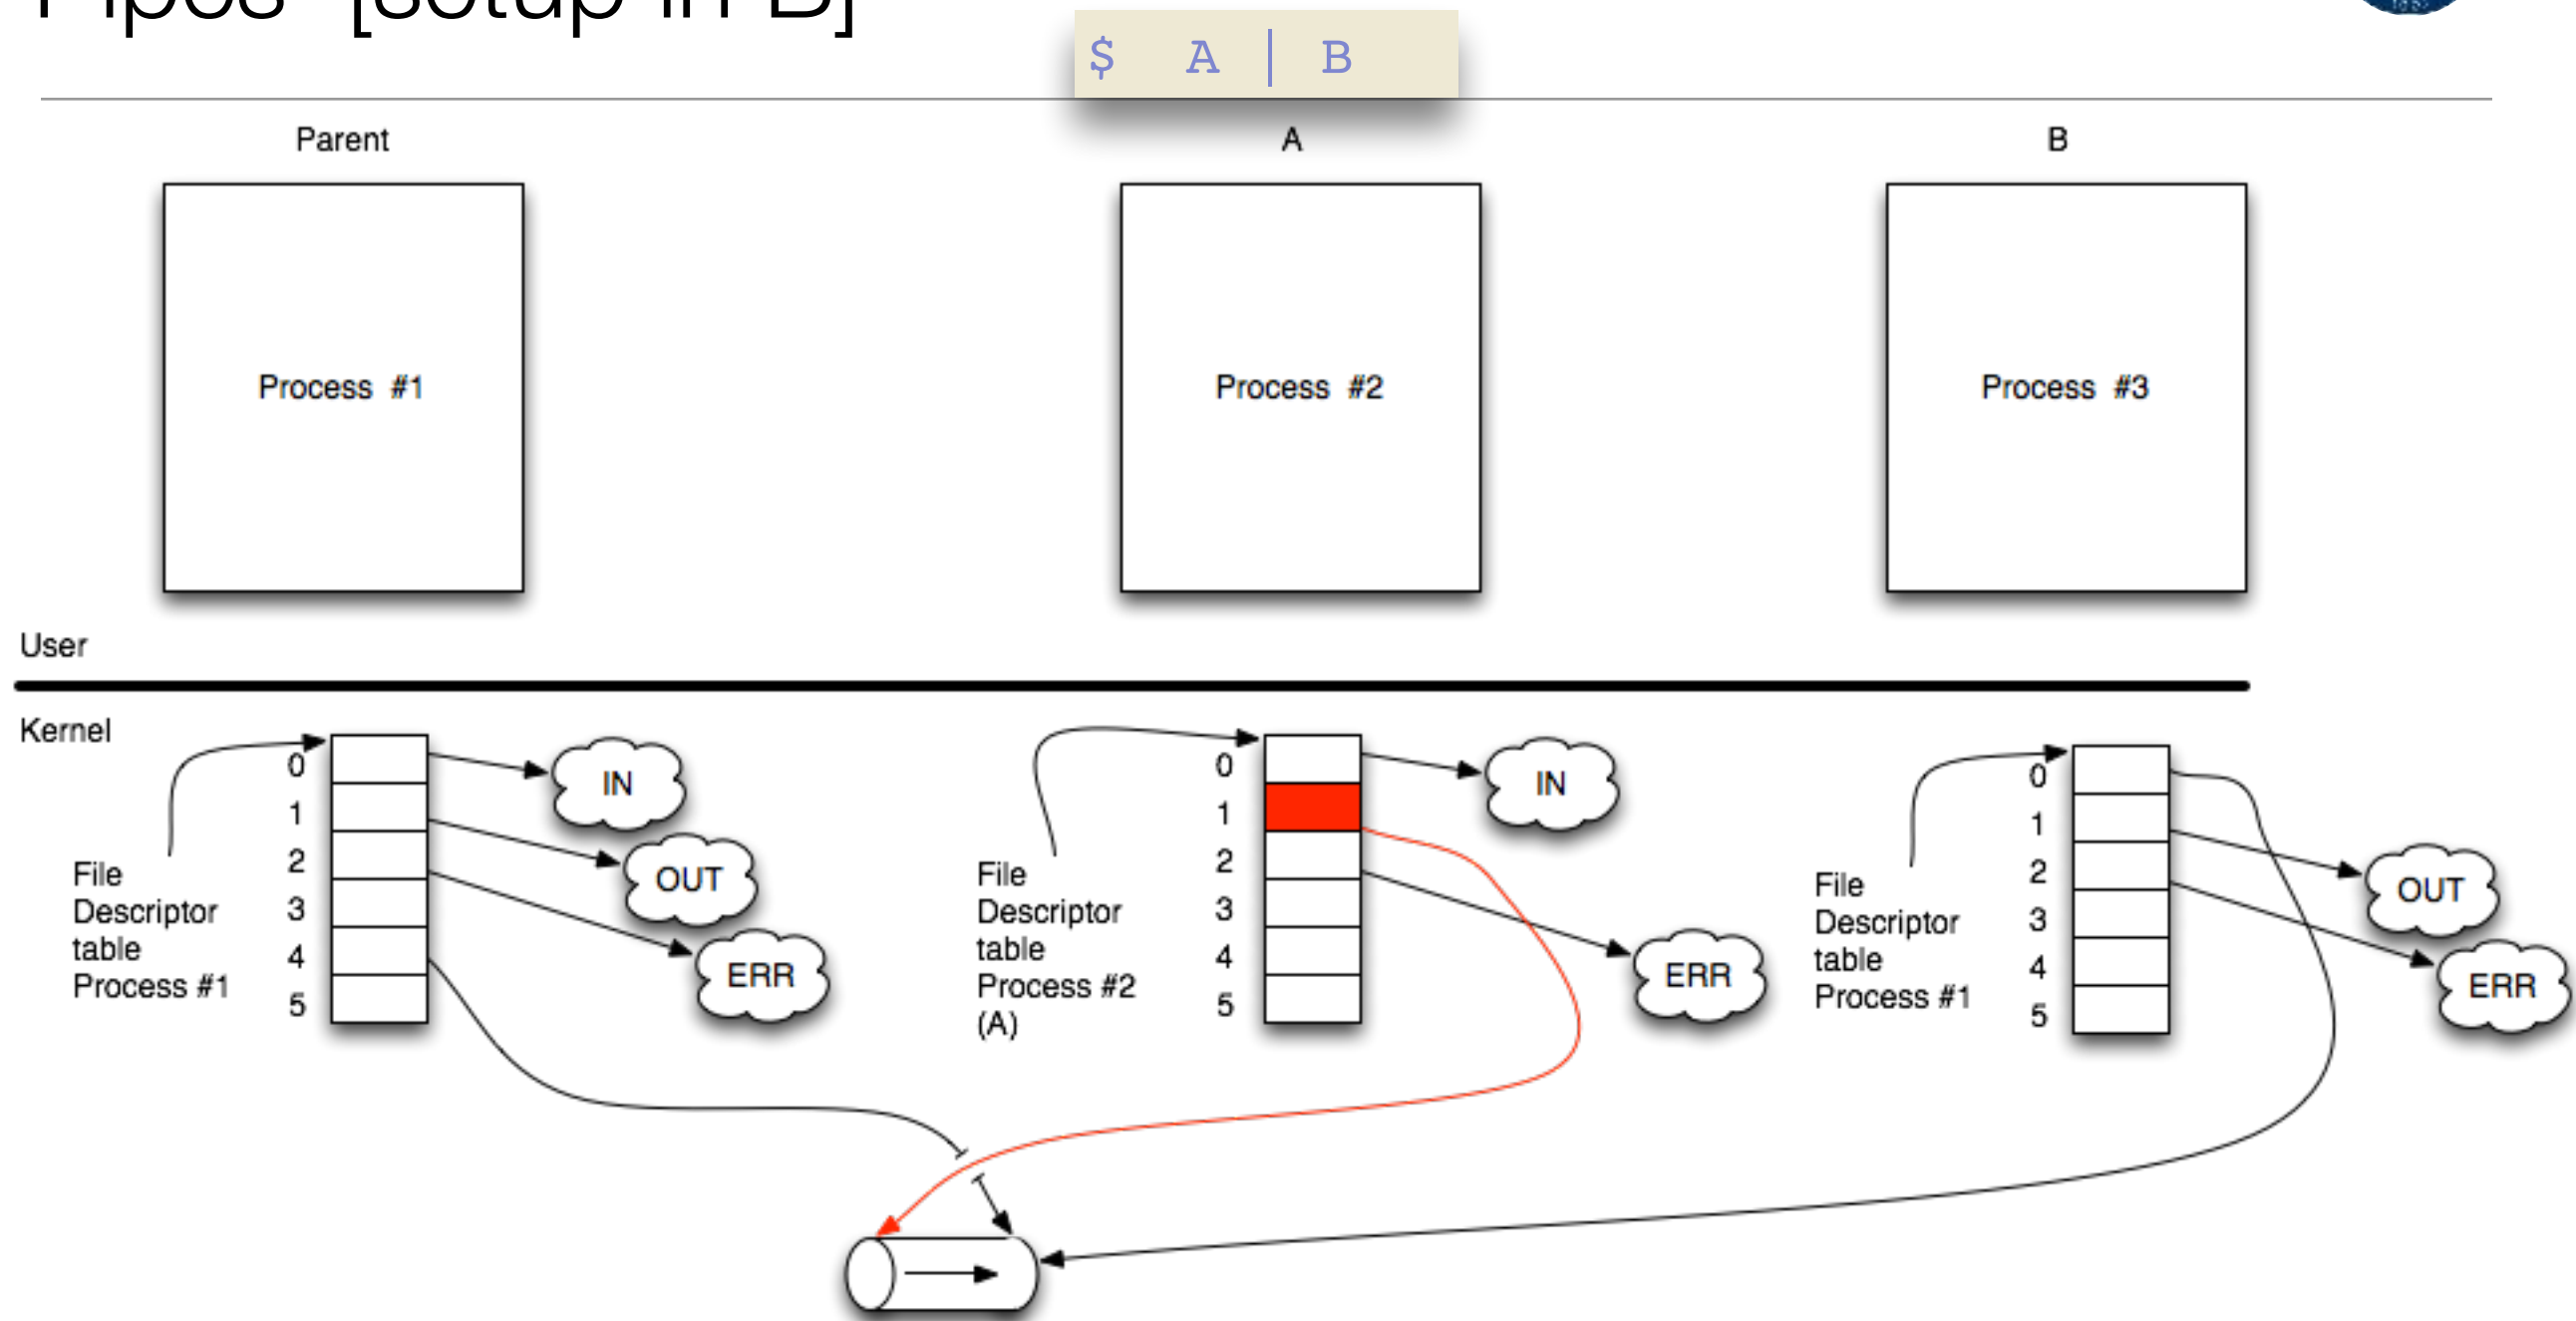


# Pipes [setup in B]





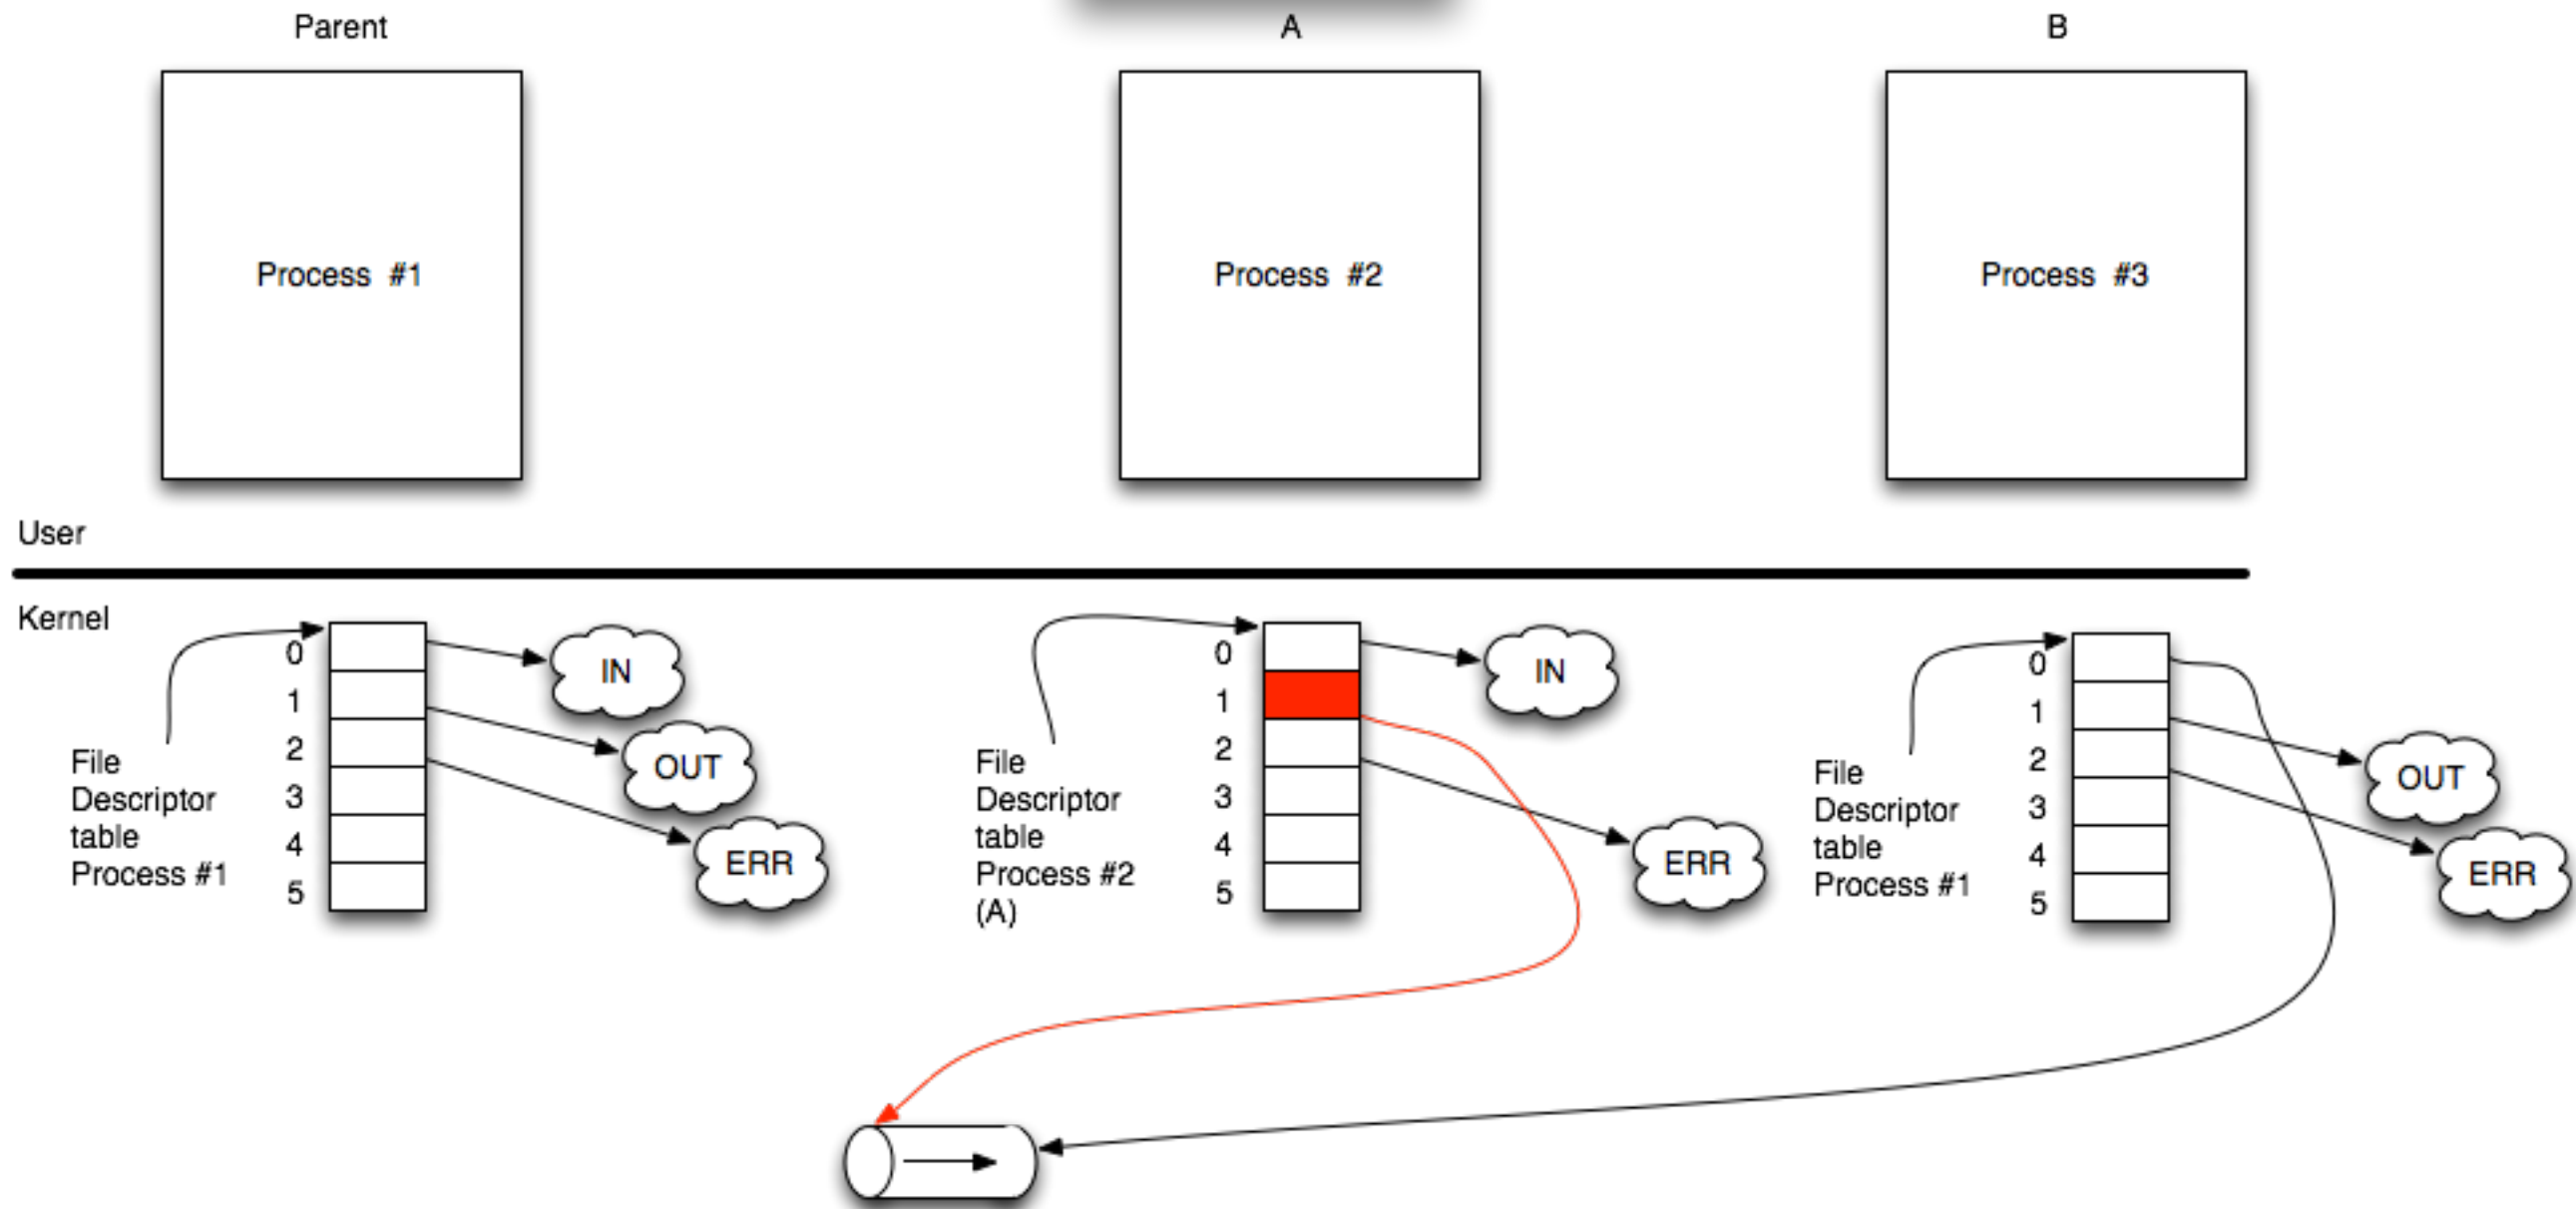
# Pipes [setup in B]





# Pipes [cleanup in Parent]

\$ A | B





# You can repeat this...

---

- Once per stage of the pipeline





# Remember

---

- **All processes in the pipeline**
  - Are running concurrently
  - As soon as data sent in the pipe...
  - The next process picks it up and starts working.
- **When the first (source) process dies...**
  - Its output is closed
  - But its child keeps reading “remaining inputs”
  - And dies when he has written its last output.



# Opened files for a dying process?

---

- **When a process ends**
  - All its opened files are automatically closed
- **Therefore....**
  - Any lingering pipes are closed automatically
- **Only danger**
  - Not closing properly in the parent (shell) process!