



## A C Primer (Part II)





# Overview

---

- Compound Types
  - Arrays
  - Structures
  - Pointers
- Pointer arithmetic
- Memory layout and alignment



# Arrays

---

- Is a type constructor
  - It produces a new type
- As expected
  - Arrays represent a linear, contiguous collection of “things”
  - Each “thing” in the array has the **same fixed** type.
- Examples
  - Array of characters
  - Array of integers
  - Array of booleans
  - Array of structures
  - Arrays of arrays....



# Array Example 1

- Simple array of integers

```
int main()  
{  
    int x[5];  
    x[0] = 1;  
    x[1] = 2;  
    x[2] = 3;  
    x[3] = 4;  
    x[4] = 5;  
    return 0;  
}
```

0	1	2	3	4
1	2	3	4	5

Arrays are 0-based  
First is @ offset 0  
Last is @ offset n-1



# A Few Questions

---

- Where is the array allocated ?
- When is the array allocated ?
- When is the array deallocated ?
- What about recursion ?
- What happens if you try to access `x[5]` ?
- What happens if you try to access `x[-1]` ?
- What if you do not know the size at *compile time*?



# A Few Questions

---

- Where is the array allocated ?

On the stack, so it is automatic

- When is the array allocated ?

It is allocated when you enter the function

- When is the array deallocated ?

When the function returns

- What about recursion ?

Each invocation gets its own copy! ;-)

- What happens if you try to access  $x[5]$  ?

Ooooooh..... you are accessing memory that is not yours!

- What happens if you try to access  $x[-1]$  ?

Same as above!

- What if you do not know the size at *compile time*?

We can deal with this later on .... [dynamic allocation]





# Array Example 2

- A character array

```
#include <stdio.h>
int main()
{
    char s[6];
    s[0] = 'H';
    s[1] = 'e';
    s[2] = 'l';
    s[3] = 'l';
    s[4] = 'o';
    s[5] = '\0';
}
```

0	1	2	3	4	5
H	e	l	l	o	0

Same as int array!  
First is @ offset 0  
Last is @ offset n-1  
Last character is \0



# Array Example 2

- A character array

```
#include <stdio.h>
int main()
{
    char s[6];
    s[0] = 'H';
    s[1] = 'e';
    s[2] = 'l';
    s[3] = 'l';
    s[4] = 'o';
    s[5] = '\0';
    printf("Array is: %s\n", s);
}
```

0	1	2	3	4	5
H	e	l	l	o	0

A 0-terminated  
character array is...

**A String!**





# Array Example 2

- A character array. Convenience initialization.

```
#include <stdio.h>
int main()
{
    char s[6] = {'H', 'e', 'l', 'l', 'o', '\0'};
    printf("Array is: %s\n", s);
}
```

0	1	2	3	4	5
H	e	l	l	o	0

You can provide an initializer list...

It works for all types



## Array Example 2

- A character array. Convenience initialization **again!**

```
#include <stdio.h>
int main()
{
    char s[6] = "Hello";
    printf("Array is: %s\n", s);
}
```

0	1	2	3	4	5
H	e	l	l	o	0

Even better for **Strings...**

You can give the list of characters  
in a double-quoted literal



## Array Example 2

- A character array. Convenience initialization **again!**

```
#include <stdio.h>
int main()
{
    char s[] = "Hello";
    printf("Array is: %s\n", s);
}
```

0	1	2	3	4	5
H	e	l	l	o	0

Even better for **Strings...**

You can even drop the size.  
C will compute it from the initializer  
(also true for int arrays)



# Array Indexing

---

- Works as in Java
  - You can read or write anywhere inside the array
- Unlike Java
  - You can also read or write *outside the array*. Be very careful!
- Indexing
  - Must evaluate to an integer
  - Can be an expression
- Each definition yields a type




# Array Assignment

- Unlike Java

- You **cannot** assign a whole array at once to another array
- Even when the types match

- Note [again]

- In Java, arrays are heap allocated
- In C, arrays can be stack allocated



```
int main() {  
    int x[10];  
    int y[20];  
    int z[10];  
    x = y;  
    x = z;  
}
```

```
src (master) $ cc char3.c  
char3.c:5:6: error: array type 'int [10]' is not assignable  
    x = y;  
    ~ ^  
char3.c:6:6: error: array type 'int [10]' is not assignable  
    x = z;  
    ~ ^  
2 errors generated.
```




# Arrays as Automatic Variables

## • Two Key Facts

- You can declare arrays inside *any function*
- The size of your array can depend on *function arguments* **[dynamic!]**

```
int silly(int n,int k) {  
    int x[n];  
    for(int i = 0;i < n;i++)  
        x[i] = 0;  
    x[0] = 0;  
    x[1] = 1;  
    for(int i = 2;i < n;i++)  
        x[i] = x[i-1] + x[i-2];  
    if (0 <= k && k < n)  
        return x[k];  
    else return -1;  
}
```

```
int main() {  
    int z = silly(10,5);  
}
```







# Automatic Array Summary

---

- **Local Arrays**

- Allocated when entering the function [automatic]
- Deallocated when leaving the function [automatic]
- **NOT** initialized
- Exist directly on the stack like other variables.

- **Size**

- Can be static [a constant]
- Can be dynamic [an argument to the function]
- Cannot be too big since it is on the stack!



# Arrays as Arguments to functions

---

- Arrays can be passed to functions!
  - With one big caveat...
- Calling convention in C
  - BY VALUE for everything....
  - EXCEPT arrays...
- Arrays are always passed as “pointers”.
  - We have to look at pointers soon!

# Arrays as Arguments to functions (Statically sized)




```
int silly(int k,int x[10]) {
    for(int i = 0;i < 10;i++)
        x[i] = 0;
    x[0] = 0;
    x[1] = 1;
    for(int i = 2;i < 10;i++)
        x[i] = x[i-1] + x[i-2];
    if (0 <= k && k < 10)
        return x[k];
    else return -1;
}
```

```
int main() {
    int t[10];
    int a0 = silly(5,t);
    int a1 = t[8];
    printf("got 2 values: %d :: %d\n",a0,a1);
    return 0;
}
```

It looks like a declaration  
array passed by “address”  
modifications are visible!

# Pictorially





# Arrays as Arguments to functions (dynamically)

```
int silly(int n,int k,int x[n]) {  
    for(int i = 0;i < n;i++)  
        x[i] = 0;  
    x[0] = 0;  
    x[1] = 1;  
    for(int i = 2;i < n;i++)  
        x[i] = x[i-1] + x[i-2];  
    if (0 <= k && k < n)  
        return x[k];  
    else return -1;  
}
```

```
int main() {  
    int nb = 10;  
    int t[nb];  
    int a0 = silly(10,5,t);  
    int a1 = t[8];  
    printf("got 2 values: %d :: %d\n",a0,a1);  
    return 0;  
}
```

Very similar  
array size is known though!  
Still passed by “address”  
modification still visible!



# Arrays? Are we done?

---

- **Nope...**
  - More on arrays once we start working with pointers.





# Compound Type

---

- Is a *type constructor*
  - It builds a brand new type



# Compound Types

- Also known as “structures”

```
#include <stdlib.h>
```

```
struct Person {  
    int    age;  
    char  gender;  
};
```

A structure  
declaration! (a type)

```
int main()  
{  
    struct Person p;  
  
    p.age = 44;  
    p.gender = 'M';  
    return 0;  
}
```

A structure *definition*! (a value)



# Compound Types

- Also known as “structures”

```
#include <stdlib.h>
```

```
struct Person {  
    int    age;  
    char  gender;  
};
```

A structure  
declaration! (a type)

```
int main()  
{  
    struct Person p = {44, 'M'};  
  
    return 0;  
}
```

A structure *definition*! (a value)  
and *initialization*



# Structures

---

- **Declaration**

- Structure have a type **name**
- Can have multiple fields
- Fields can have any legal type
  - Basic types
  - Structures
  - Arrays
  - [and the other types left to discover!]

- **Definition**

- Define a value.
- Value lives on the stack (automatically deallocated on return)



# Composing Types

---

- You can compose
  - Basic Types
  - Structure Types [Compounds]
  - Array Types
- A more realistic struct example



# Realistic Example

- Embed an array in the structure for the person's name
- Make an array of structures of type Person for the whole family.
- Nest initializers
- Caveats
  - Names cannot be > 32 long.
  - Four persons in family
    - Indexed 0..3

```
#include <stdlib.h>

struct Person {
    int    age;
    char   gender;
    char  name[32];
};

int main()
{
    struct Person family[4] = {
        {50, 'M', "Darth Vader"},
        {49, 'F', "Padmé"},
        {21, 'F', "Leia"},
        {19, 'M', "Luke"}
    };
    int kidAge = family[3].age;
    return 0;
}
```





# Type Definitions

- Type names can become long
- C provide the ability to define type abbreviations
  - `typedef` declaration
    - Give existing type
    - Give new type name
  - Use the new type anywhere
- Useful to make code even more readable

```
struct Person {
    int     age;
    char    gender;
    char    name[ 32 ];
};

typedef struct Person TPerson;

int main()
{
    TPerson family[4] = {
        {50, 'M', "Darth Vader"},
        {49, 'F', "Padmé"},
        {21, 'F', "Leia"},
        {19, 'M', "Luke"}
    };

    return 0;
}
```



# Overview

---

- Basic Types
- Compound Types
  - Arrays
  - Structures
  - Pointers
- Pointer arithmetic
- Memory layout and alignment



# Pointers

---

- Perhaps the scariest part of C
- Yet....
  - The most useful part of C!
- Pointer is simply....
  - A value
  - Denoting the address of a memory cell



The idea

*"Pointers are a difficult concept" is a*

**MYTH**



# Key Insight...

- Picture it...

A PICTURE IS WORTH  
A THOUSAND  
WORDS

- And you **will** understand

- the memory model

- pointers



# Pointers Usage

---

- What can you do ?
  - Get the “address” of something
  - Dereference an “address” to get to something
  - Compute the address of something





# Memory Model...

---

- Three pools of memory
  - Static
  - Stack
  - Heap
- Each pool features
  - Different lifetime
  - Different allocation/deallocation policy

**Why does it matter?**



# Static Memory Pool

---

- This is where
  - All constants are held
  - All strings in the program are held
  - All variables declared “static” are held
- Allocated when
  - The program start
- Deallocated when
  - The program terminates
- Bottom line
  - **FIXED SIZE**

Rule for this class

DO NOT USE STATIC

**NEVER, EVER, EVER**

(repeat on a T. Swift tune...)



# Automatic Memory Pool

---

- This is where....

- Memory comes from for **local variables** of functions!
- Allocated automatically when entering the function
- De-allocated automatically when you leave the function.

- Recursion ?

- Each recursive invocation gets its own set of local variables!

- Bottom line

- It's easy to manage (automatic!)
- It's variable over time
- Scope is that of function.



# Dynamic Memory Pool

---

- **This is where...**
  - Memory comes from for manual “on-the-fly” allocations
  - Two simple APIs (at the core)
    - `malloc(b)` [allocate a block of ‘b’ bytes]
    - `free(p)` [frees a previously allocated block]
- **Who is in charge ?**
  - The programmer for both allocation / deallocation
- **Lifetime of memory blocks ?**
  - As long as they are not freed!



- Taking the address of....

- A static ?

- The address is never going to go “bad”
- The static lives as long as the program!

- A stack [automatic] variable ?

- The address is valid as long as the variable is!
- When the function returns.... The address is bogus

- A heap variable ?

- The address is valid as long as the variable is!
- The variable disappear when explicitly de-allocated (freed)





# Pointer examples

- Simple deal
  - & “makes” an address
  - \* “dereferences” an address
- It works for all types

```
int x = 10;  
  
int* px = &x;  
  
*px = 20;
```





# Declaration

- Word to the wise...

- The following three are identical

```
int*    p;  
int  *  p;  
int    *p;
```

**Notation**

- They all declare...

- p to be a pointer to an integer

- But

- First one m
- Second on
- Third says

The first is very clear  
But...  
The third is classic C



# Pitfall

---

- Read the following !

```
int    *a,b;
```

- What is...

- a ?
- b ?





# How Much Space ?

---

- How do you determine the amount of space for some type?
  - You need this to dynamically allocate space!
- Easy!

```
sizeof(T)
```

- Returns the number of bytes to hold a value of type **T**



# Pointer and Dynamic Memory

---

- **Pointers are useful when**
  - Allocating memory at run time!
- **Example**
  - You need n integers (n known at runtime)....

```
void doSomething(int n)
{
    int* pox;

    pox = (int*)malloc(sizeof(int)*n);

    *pox = 0;    // What happens?

    free(pox);
}
```



# Shorthand

- If you need to allocate an array
  - There is another library function called `calloc`
  - `calloc` is implemented in term of `malloc`
  - `calloc` also initialize the content to 0



Library

```
void doSomething(int n)
{
    int* pox;

    pox = (int*)calloc(n, sizeof(int));

    *pox = 0;    // What happens?

    free(pox);
}
```



# Deallocation

---

- **Straightforward**

- Simply call the library function “free”
- Takes a pointer to the block to free



```
free(ptr)
```



# Best Practice

---

- Always remember two key rules

## **RULE #1:**

Everything you malloc should eventually be freed

## **RULE #2:**


Only free what is allocated via malloc / calloc

- Consequences of not remembering the rules
  - Memory “Leaks” [you will eventually run out of memory]
  - Double deletion and horrible crashes



# Pointers and Arrays

- What really happened....



- Wait!!!!
  - That looks like an array!



# Pointers and Arrays

- Yes, pointers and arrays are the same thing
  - Array is represented by the address of 0th element!


```
void doSomething(int n) {  
    int* pox;  
  
    pox = (int*)malloc(sizeof(int)*n);  
  
    /*pox = 0;    // What happens?  
  
    pox[0] = 0;  
    pox[1] = 1;  
    ...  
    pox[n-1] = n-1;  
    free(pox);  
}
```



# Pointer Arithmetic

- Since pointers and arrays are the same thing....

- What is going on when we write: `pox[3] = ...;`



- So `pox[3]` is
- 12 bytes into the `pox` “array”
  - 3rd element of `pox` “array”






# Pointer Arithmetic

- Since pointers and arrays are the same thing....

- What is going on when we write:

```
pox[3] = ...;
```



So `pox[3]` is

- 12 bytes into the `pox` “array”
- 3rd element of `pox` “array”

```
*(pox+3) = ...; // arithmetic is type driven
```



# Implication

---

- You can compute the address of anything
- You only need to
  - Know the base address
  - Know the offset
  - Know the types



# Bounds

- When doing...

```
int n = 6;  
int* pox;  
pox = (int*)malloc(sizeof(int)*n);
```

- Index of

- First element is 0
- Last element is  $n-1$  ( $6-1 = 5$ )

- Yet...

- C allows reads and write
  - Before 0
  - After  $n-1$





# Initializing Pointers

---

- Ways to initialize a pointer

- As a result of malloc

```
int *p = (int*)malloc(sizeof(int)*10);
```

- As a result of &

```
int a = 10;  
int *p = &a;
```

- As a result of assigning NULL

```
int* t = NULL;
```



# Best Practice

---

- You should know that....
  - A call to `malloc` may **fail**
    - Whenever you are out of virtual memory
    - You would get back the NULL value
    - Not much to do except report the error and terminate *nicely*.
- Idiom

```
int* p;  
p = (int*)malloc(sizeof(int)* n);  
if (p == NULL)  
    report error and finish;
```



# Back to the **scanf** example!

```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    printf("%s ate %d apples and %d pears.\n", name, apples, pears);
    return 0;
}
```

- What is going on ?
  - Pass to scanf three **VALUES**
    - name
    - **address of** pears     **[&]**
    - **address of** apples     **[&]**




# In Picture (Refresher!)

- **Memory....**

- Every *Process* has an
  - **Address Space**
- **Executable** code is at the bottom
- **Statics** are just above
- **Stack** is at the top (going down!)
- **Heap** grows from the bottom (going up!)
- Gray no-man's land is up for grab

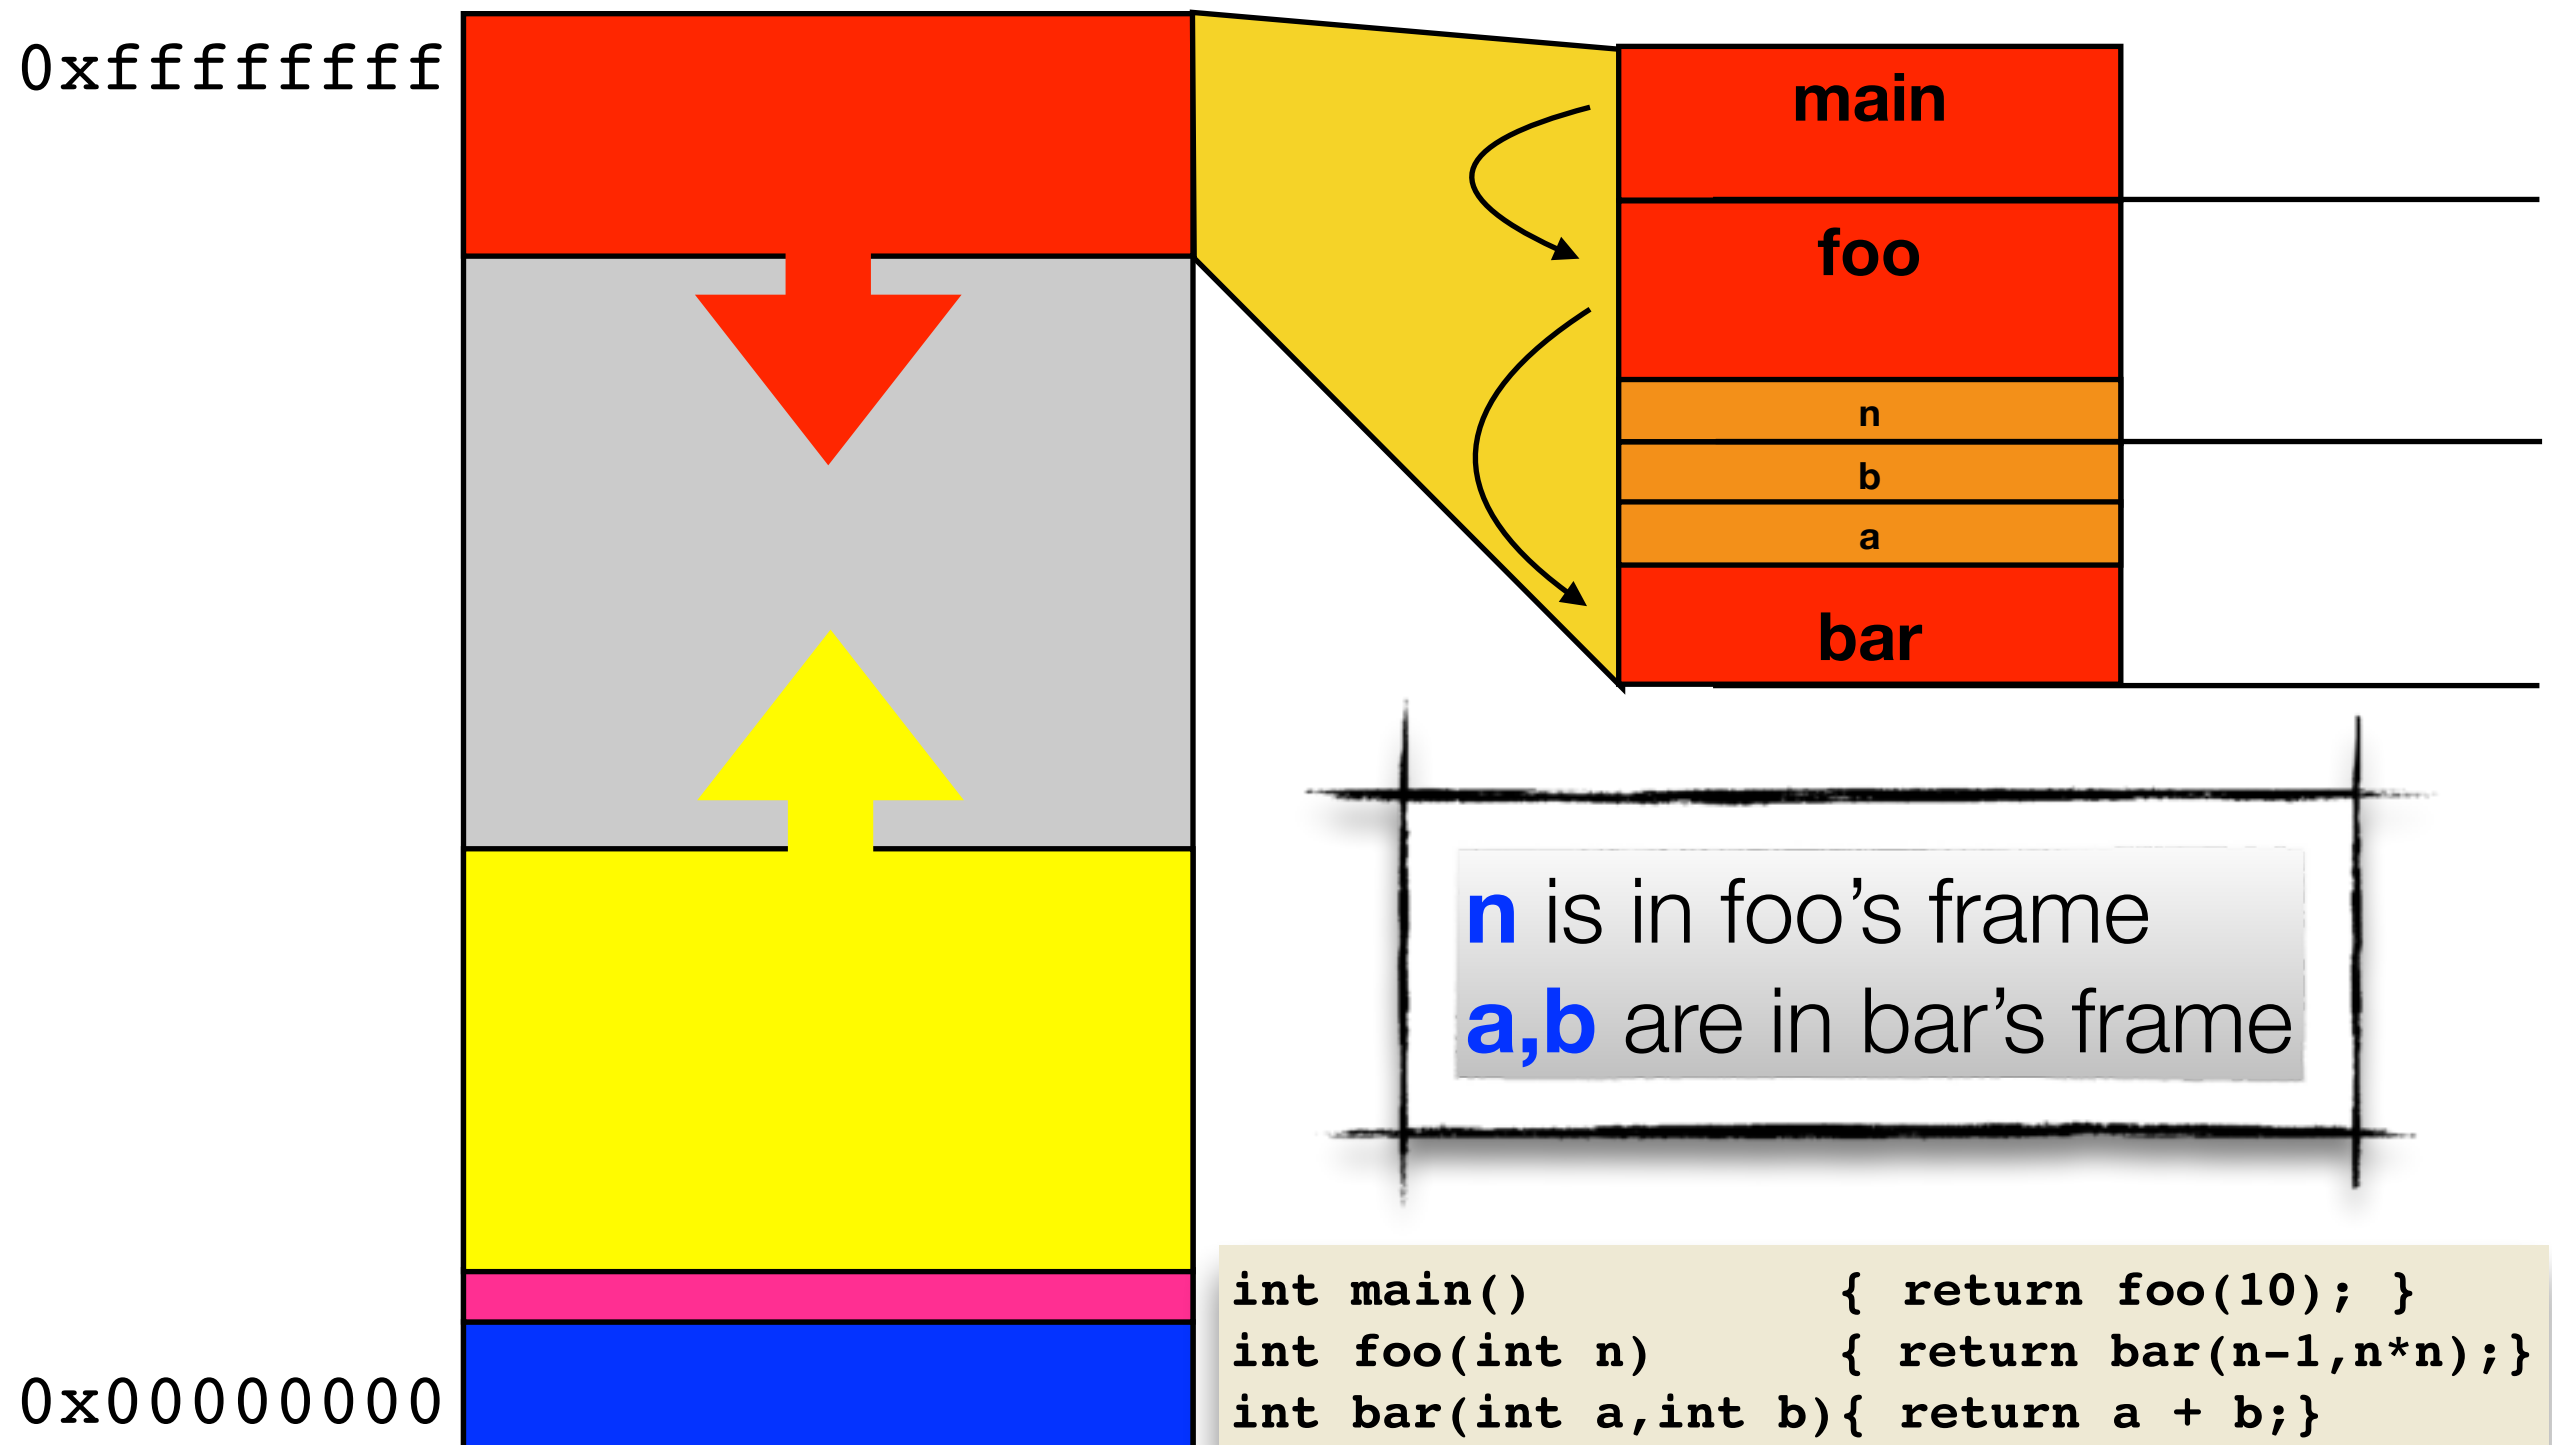
High 0xfffffffffff

LOW 0x00000000





# Zooming in [on the stack!]







# scanf Code

- The symbols are coming from ?

- name ?
- pears ?
- apples ?

- What are their **types** ?


- name ?
- pears ?
- apples ?

```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Frame of main ... In picture




```
#include <stdio.h>

int main()
{
    char name[128];
    int pears  = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Calling **scanf** per se




```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Calling **scanf** per se

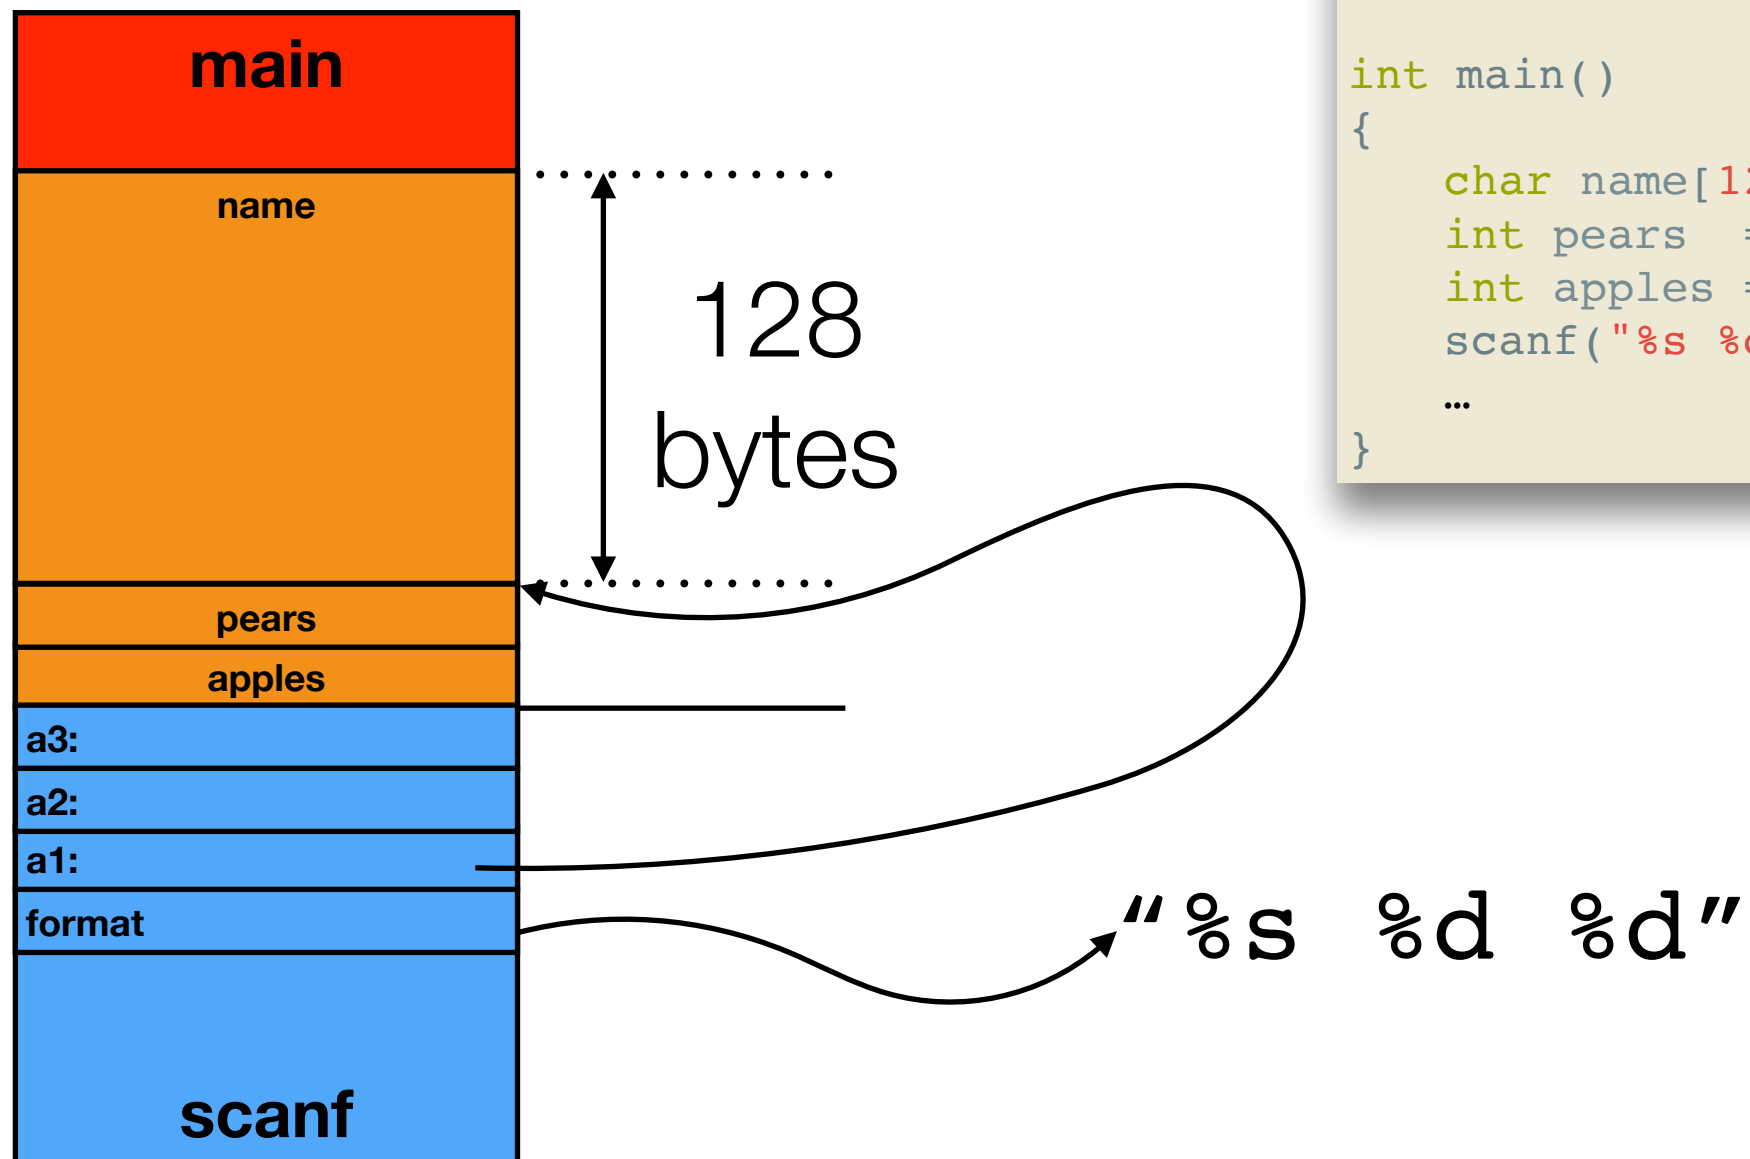


```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Calling **scanf** per se

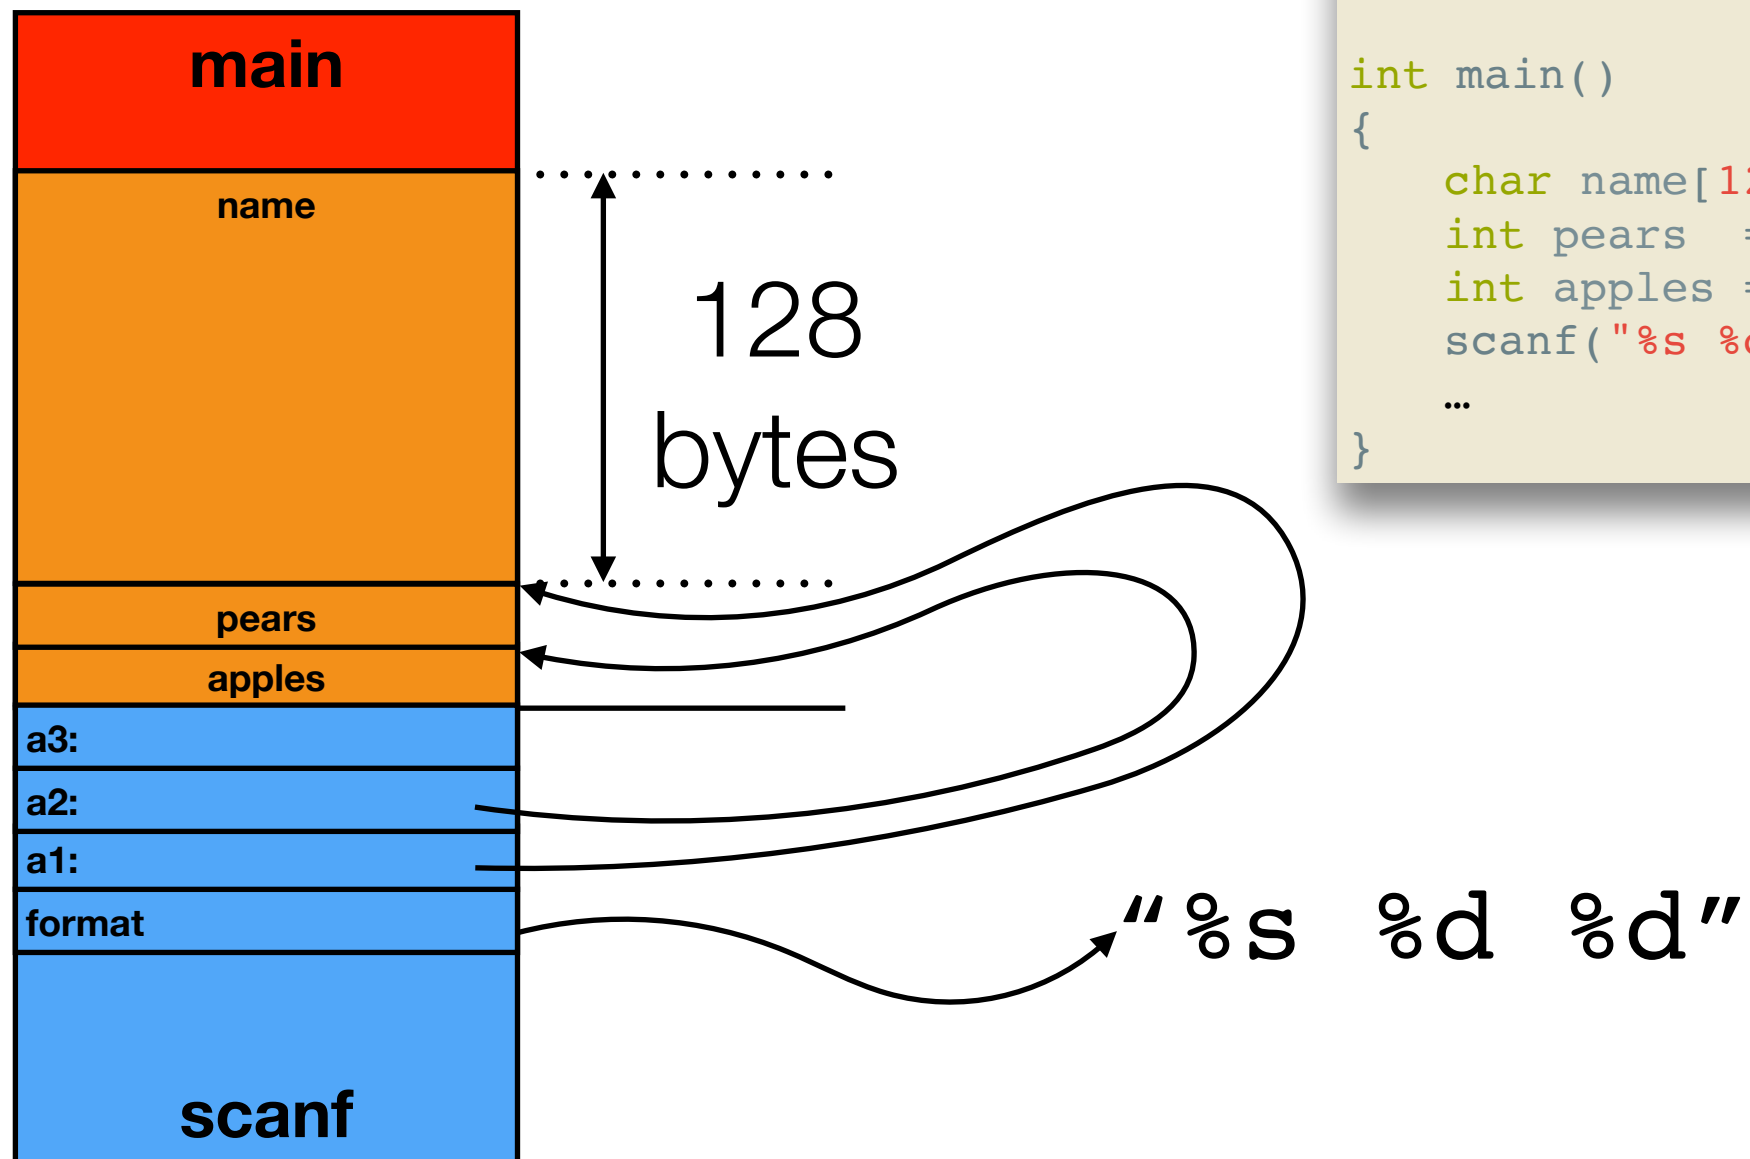


```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Calling **scanf** per se




```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```



# Calling **scanf** per se



```
#include <stdio.h>


int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```

"%s %d %d"

We need **&** on pears / apples to give their addresses to **scanf** as **scanf** must **WRITE** there!



# Calling **scanf** per se



```
#include <stdio.h>

int main()
{
    char name[128];
    int pears = 0;
    int apples = 0;
    scanf("%s %d %d", name, &pears, &apples);
    ...
}
```

"%s %d %d"

We **do not** need **&** on name as name is an array and therefore it already is a pointer!  
Thus, scanf can **WRITE** there too!





# Overview

---

- Basic Types
- Compound Types
  - Arrays
  - Structures
  - Pointers
- Pointer arithmetic
- Memory layout and alignment



# Pointers are addresses

---

- It's quite simple

- The **value** of a pointer is the **address** in the address space
- Thus, the value of a pointer is an integer between (in hexa, 32-bit)
  - 0x00000000 [on 64-bit: 0x0000000000000000]
  - 0xFFFFFFFF [on 64-bit: 0xFFFFFFFFFFFFFFFF]

- Corollary

- If a pointer is a integer, you can do *arithmetic...*
- To *compute* other addresses



# Rules of Pointer Arithmetic

---

- Very similar to integer arithmetic
- You can
  - Add a value to a pointer
  - Subtract two pointers
  - Scale an offset
- Big difference
  - All additions are subject to **automatic scaling** of the constant added

Main Driver in **automatic scaling** is  
The **type** of what is pointed to




# Example

- Simple illustration

```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3]    = 10;
    int *p    = tab + 3;
    printf("What is at tab+3? = %d\n", *p);
    *p = 20;
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```






# Example

- Simple illustration

```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3]    = 10;
    int *p    = tab + 3;
    printf("What is at tab+3? = %d\n", *p);
    *p = 20;
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```






# Example

- Simple illustration

```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3]    = 10;
    int *p    = tab + 3;
    printf("What is at tab+3? = %d\n", *p);
    *p = 20;
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```






# Example

- Simple illustration

```
#include <stdlib.h>


int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3]    = 10;
    int *p    = tab + 3;
    printf("What is at tab+3? = %d\n", *p);
    *p = 20;
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```





# But what about memory addresses?

- Same story...!



tab+1 = ?


# 0x11110001 ?





# But what about memory addresses?

- Same story...!



0x11110000

tab+1 = ?

~~0x11110001 ?~~

0x11110004 ?

## WHY?

Simply because tab is a pointer to an int and an int is 4-bytes wide!



# Bottom line

```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3]    = 10;
    int *p    = tab + 3;
    printf("What is at tab+3? = %d\n",*p);
    *p = 20;
    printf("What is at tab[3]? = %d\n",tab[3]);
    return 0;
}
```

The offset **3** is scaled by the compiler with the size of the type to get an address in bytes





# Effect of casting types ?

---

- If you cast a pointer type...
  - Any subsequent pointer arithmetic will use the type you chose
  - Hence the scaling constants are different
- Note
  - `sizeof(char) = 1`
- Corollary
  - Casting a pointer to (**char\***) results in no scaling!

This is why you need to cast the result from a call to malloc!



# Fundamental observation

---

- **C is permissive**

- It will let you cast pointers in any way you like
- This is called “weak typing”

- **Corollary**

- You can “forge” pointers to point wherever you wish....
- [within your address space of course!]

That's what makes C very attractive for  
low-level programming

That is also very **powerful** and thus **dangerous!**



# Pointer subtraction?

---

- **Very easy to understand!**
  - It computes the offset (distance) between the two pointers!
  - The distance is in number of “Data Items” (again, based on types)
- **Requirement**
  - Both pointers have the same type
- **Check the example!**



# Example

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    int* t = (int*)malloc(sizeof(int)*10);
    int* p = t + 9;
    int dist = p - t;
    printf("Distance is %d\n",dist);
    return 0;
}
```


## Output

```
src (master) $ cc ptrsub.c
src (master) $ ./a.out
Distance is 9
src (master) $
```



# Pointer comparisons

- You can also compare two pointers
  - `<` , `>` , `<=` , `>=` , `!=` , `==`
- Purpose
  - Check boundary conditions in arrays
  - Manually manage memory blocks
- Semantics
  - Simply based on process layout!
    - `==` , `!=` are obvious
    - `<` , `>` , `<=` , `>=` easy enough!







# Generic Pointers

---

- **Purpose**

- Have a pointer to a memory block whose content is “un-typed”
- Very convenient for raw memory manipulation

- **Requires**

- Casting the pointer types before dereferencing for read / write



# Generic Pointers Example

- **Note**

- Arithmetic is possible
- Arithmetic is “byte-oriented”

```
#include <stdio.h>
#include <stdlib.h>
int main()
{
    void* p = (void*)malloc(sizeof(int)*10);
    void* a = p + 4;
    printf("p is %p  a is %p. Distance is: %ld\n",p,a,a - p);
    return 0;
}
```

```
src (master) $ cc void.c
src (master) $ ./a.out
p is 0x7fbe92c03a20  a is 0x7fbe92c03a24. Distance is: 4
```



# Overview

---

- Basic Types
- Compound Types
  - Arrays
  - Structures
  - Pointers
- Pointer arithmetic
- Memory layout and alignment



# C Structures

---

- Structures are a way to **package** related attributes

```
struct IntStack {  
    int tab[32];  
    int top;  
};  
  
...  
struct IntStack myStack;  
myStack.top = 0;
```

- Fields are **contiguous** in memory
- Fields are **aligned** for the natural types

# Alignment

---



Alignment requirement on a x64 architecture	
char	1
short	2
int	4
long	8
float	4
double	8



# Meaning...

---

- Memory used to store a value of type X **MUST**
  - be lined-up on a multiple of natural alignment for X
- Why?
  - Performance!
- If you do not respect alignment requirements...
  - BUS ERROR (sigbus)
  - The O.S. will kill your program

# Good news...

---



The C compiler handles alignment  
automatically  
99% of the time

# Bad news...

---




The C compiler handles alignment  
automatically

**99%** of the time

The Programmer must handle the remaining 1%



# Example




```
struct {  
    char x[5];  
    int y;  
    double z;  
}
```

Can you do anything about  
this to *improve*?



# Improvement

Padding to multiple of 4



```
struct {  
    char x[5];  
    int y;  
    double z;  
}
```

```
struct {  
    double z;  
    int y;  
    char x[5];  
}
```



# Padding ?

---

- **Motivation....**

- What if you wish to create an array of structs ?
- Second struct in array should start on 4 boundary!

- **Conclusion**

- Make it easy
- All structures have sizes that are multiple of
  - 4 on i32
  - 8 on x64



# What is this 1% business?

---

- **When you do pointer arithmetic of course!**
  - You must understand the layout (and padding) of structures to correctly compute addresses within the structure!
- **When you call system routines with specific alignment needs**
  - Your arguments must comply
  - Use compiler annotations to force specific alignments (beyond our scope, simply remember that this exists!)