

Memory Operations



Caiwen Ding

Department of Computer Science and Engineering
University of Connecticut

CSE3666: Introduction to Computer Architecture

Outline

- Memory
- Load/store instructions
 - Move data between registers and memory
- Data of other types
 - Words, halfwords, and bytes
 - ASCII strings
- Address alignment
- Endianness

Reading: Sections 2.3.

References: Reference card in the book.

Memory

- Memory is an array of bytes
- Each byte is numbered. The number is the **address**
- Each address identifies a byte
 - If a data item is larger than one byte, its address is the first byte in memory
- A 32-bit address space supports 4 GiB
 - A 64-bit address space supports 16 EiB (exbibytes)



Kibibytes (KiB) vs kilobytes (KB)

- We always mean KiB, MiB, GiB

Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	10^3	kibibyte	KiB	2^{10}	2%
megabyte	MB	10^6	mebibyte	MiB	2^{20}	5%
gigabyte	GB	10^9	gibibyte	GiB	2^{30}	7%
terabyte	TB	10^{12}	tebibyte	TiB	2^{40}	10%
petabyte	PB	10^{15}	pebibyte	PiB	2^{50}	13%
exabyte	EB	10^{18}	exbibyte	EiB	2^{60}	15%
zettabyte	ZB	10^{21}	zebibyte	ZiB	2^{70}	18%
yottabyte	YB	10^{24}	yobibyte	YiB	2^{80}	21%

A video on Kilobyte or Kibibyte?

<https://www.youtube.com/watch?v=ZRQVPcgf5yE>

Using data in memory

- Many ISAs like RISC-V **cannot** compute on data in memory directly
 - Must load data into a register first
- Two kinds of instructions to exchange data between registers and memory
 - **Load** : memory to register
 - **Store** : register to memory
- Need to know the **address** to read/write memory
 - You need an address to save/fetch items

Variables defined in your program

```
.align 2
# a word with initial value 3
x:      .word 3

# two words with initial values
y:      .word 4, 5
```

```
// in C
int  x = 3;
int  y[2] = {4, 5};
```

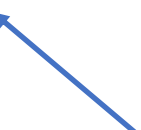
How do you get the address of
a variable in a register?

	Address	Value
	0x00FE 901C	
	0x00FE 9018	
	0x00FE 9014	5
y	0x00FE 9010	4
x	0x00FE 900C	3
	0x00FE 9008	
	0x00FE 9004	
	0x00FE 9000	

How to get the address of a variable in a register?

- Basically, we need to load a 32-bit constant in a register
- We can also use a pseudoinstruction **LA**
 - It is converted into a couple of real instructions
 - We will learn the real instructions later

la s1, x



a label in assembly
or
a variable name

Load/Store instructions

```
# load a word from mem into rd  
# Reg[rd] = Mem[Reg[rs1] + offset]
```

```
lw    rd,    offset(rs1)
```

```
# save a word to mem  
# Mem[Reg[rs1] + offset] = Reg[rs2]  
sw    rs2,    offset(rs1)
```

- Load/store words
- Offset is also called displacement
- The **effective address** is the value in register plus the offset

$$\text{address} = \text{Reg[rs1]} + \text{offset}$$

Effective address is the
calculated memory address

Write 0(rs1) even if the offset is 0
Assembler may support "(s1)" as a pseudoinstruction

Example

- Each row in the table is a byte
- Assume a's address is in s1. Write RISC-V instructions to do

```
int a, b;
```

```
b = a
```

Var name	Address	Value
b	0x00FE 9007	
	0x00FE 9006	
	0x00FE 9005	
	0x00FE 9004	
	0x00FE 9003	
a	0x00FE 9002	
	0x00FE 9001	
	0x00FE 9000	

Answer

- Each row in the table is a byte
- Assume a's address is in s1. Write RISC-V instructions to do

`b = a;`

```
lw    t0,0(s1)
sw    t0,4(s1)
```

Var name	Address	Value
b	0x00FE 9007	
	0x00FE 9006	
	0x00FE 9005	
	0x00FE 9004	
	0x00FE 9003	
a	0x00FE 9002	
	0x00FE 9001	
	0x00FE 9000	

Array

Suppose word array A starts from 0x9000 (stored in s1).

How do we read/write A[0], a[1], etc.?

Given a valid index i, how do we access A[i]?

C	Address	Value
	...	
A[7]	0x901C	700
A[6]	0x9018	600
A[5]	0x9014	500
A[4]	0x9010	400
A[3]	0x900C	300
A[2]	0x9008	200
A[1]	0x9004	100
A[0]	0x9000	0

4 bytes in A[1]

Address	Value
0x9007	0x00
0x9006	0x00
0x9005	0x00
0x9004	0x64

Memory Example

C code:

```
A[20] = h + A[5];
```

A is a word array.

Variable	Register
h	s2
A's addr	s3

Memory Example

C code:

```
A[20] = h + A[5];
```

Variable	Register
h	s2
A's addr	s3

A is a word array.

RISC-V code

```
lw    t0, 20(s3)    # load A[5]
add   t1, t0, s2
sw    t1, 80(s3)    # save to A[20]
```



Example: Clearing an array

```
// assume a's address is in s1
for (i = 0; i < 8; i = i + 1)
    a[i] = 0;
```

How do we do this ???



Address	Value
0x9024	
0x9020	
0x901C	a[7]
0x9018	a[6]
0x9014	a[5]
0x9010	a[4]
0x900C	a[3]
0x9008	a[2]
0x9004	a[1]
0x9000	a[0]

Clearing an array - v1

- **Array indexing** involves
 - Multiplying index by element size
 - Adding to array base address

```
for (i = 0; i < 8; i = i + 1)
```

```
    a[i] = 0;
```

```
    i = 0
```

```
    goto test
```

```
loop:
```

```
    Compute 4*i
```

```
    Add to base address (in s1)
```

```
    Write to the address
```

```
    Increment i
```

```
test: If (i < 8) goto loop
```

Address	Value
0x9024	
0x9020	
0x901C	a[7]
0x9018	a[6]
0x9014	a[5]
0x9010	a[4]
0x900C	a[3]
0x9008	a[2]
0x9004	a[1]
0x9000	a[0]

Example: array copying

C code:

```
for (i = 0; i < 100; i ++)  
    B[i] = A[i];
```

Variable	Register
i	s1
A's addr	s2
B's addr	s3

A and B are word arrays.

```
for (i = 0; i < 100; i ++){  
    t = A[i];    # how do we do this ???  
    B[i] = t;  
}
```


Array copying - v1

copy array. array version

```
for (i = 0; i < 100; i ++)  
    B[i] = A[i];
```

Variable	Register
i	s1
A's addr	s2
B's addr	s3

RISC-V code

```
        li      s4, 100  
        li      s1, 0  
        beq     x0, x0, test # we know s1 < s4  
loop:  
        slli    t0, s1, 2      # t0 = i * 4  
        add     t2, t0, s2     # compute addr of A[i]  
        lw      t1, 0(t2)  
        add     t3, t0, s3     # compute addr of B[i]  
        sw      t1, 0(t3)  
        addi    s1, s1, 1  
test:   bne     s1, s4, loop # 7 instructions in the loop
```

Address alignment

- Alignment: Data item's address is a multiple of its size
 - Address of words is a multiple of 4
 - Address of half words is a multiple of 2
- Data addresses do not have to be aligned in RISC-V, but **misalignment will cause poor performance**
 - The addresses must be aligned in this course!

```
# align the address of next variable to  $2^2 = 4$   
    .align 2
```

You want to sit with you family when you fly!

Byte order

How is a word stored in memory?

x1 is 0x01020304

sw x1, 0x100(x0)

Which byte goes to address 0x100?

Memory Address	Value
0x0000 0103	
0x0000 0102	
0x0000 0101	
0x0000 0100	

Endianness

```
# x1 is 0x01020304
```

```
sw    x1, 0x100(x0)
```

Big-endian: The **highest** byte goes to the lowest memory address.

Little-endian: The **lowest** byte goes to the lowest memory address.

Memory Address	Value
0x0000 0103	04
0x0000 0102	03
0x0000 0101	02
0x0000 0100	01

Memory Address	Value
0x0000 0103	01
0x0000 0102	02
0x0000 0101	03
0x0000 0100	04



RISC-V uses little endian.

Question

What are the bits in t0 after the following instruction?

lw t0, 0x200(x0)

- A. 0x3265 81AC
- B. 0xAC81 6532
- C. 0xCA18 5623
- D. 0x6532 AC81
- E. None of the above

Memory Address	Value
0x0000 0203	0x32
0x0000 0202	0x65
0x0000 0201	0x81
0x0000 0200	0xAC

Data of other sizes

- RISC-V supports data of other sizes
 - Each type can be signed or unsigned

Number of bits	Name	C types (typical)
8 bits	byte	char
16 bits	half word	short int
32 bits	word	int, long int

```
# load signed (sign extended) byte/halfword  
lb/lh    rd, offset(rs1)
```

```
# load unsigned (0 extended) byte/halfword  
lbu/lhu  rd, offset(rs1)
```

```
# Store the lowest byte/halfword  
sb/sh    rs2, offset(rs1)
```

Load/store instructions

addr is the same for all load/store instructions

addr: `offset(rs1)`

Data size	Load signed	Load unsigned	Store
Word (32 bits)	<code>lw rd,addr</code>		<code>sw rs2,addr</code>
Half word (16 bits)	<code>lh rd,addr</code>	<code>lhu rd,addr</code>	<code>sh rs2,addr</code>
Byte (8 bits)	<code>lb rd,addr</code>	<code>lbu rd,addr</code>	<code>sb rs2,addr</code>

Why is there no 'lwu' here?

Strings in our programs

We will only deal with ASCII strings in this course

```
s:      .string      "CSE3666"      # or use .asciz
```

print a string (terminated by null)

```
la      a0, s
```

```
addi    a7, x0, 4
```

```
ecall
```

// in C

```
char    s[10] = "CSE3666";
```

What is the value in a0 after la?

Address	Value
0x00FE 9017	0
0x00FE 9016	54
0x00FE 9015	54
0x00FE 9014	54
0x00FE 9003	51
0x00FE 9002	69
0x00FE 9001	83
s 0x00FE 9000	67

Example: string copy

Copy string s to d.

Use pointers.

Variable	Register
s's addr	a1
d's addr	a0
c	t0

// C code

```
char c;
```

```
do {
```

```
    c = *s;
```

```
    *d = c;
```

```
    s += 1;
```

```
    d += 1;
```

```
} while (c);
```

*s means s[0]

*d means d[0]

Example: string copy answer

Copy string s to d.

Variable	Register
s's addr	a1
d's addr	a0
c	t0

// C code

```
char c;  
do {  
    c = *s;  
    *d = c;  
    s += 1;  
    d += 1;  
} while (c);
```

RISC-V

```
loop:  
    lb    t0, 0(a1)  
    sb    t0, 0(a0)  
    addi  a1, a1, 1  
    addi  a0, a0, 1  
    bne   t0, x0, loop
```

Question

What is the value in t0 after the following instruction?

```
lb    t0, 0x201(x0)
```

- A. 0x0000 00AC
- B. 0x0000 0081
- C. 0xFFFF FFAC
- D. 0xFFFF FF81
- E. None of the above

Memory Address	Value
0x0000 0203	0x32
0x0000 0202	0x65
0x0000 0201	0x81
0x0000 0200	0xAC

Registers vs. Memory

- Registers are faster to access than memory
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler must use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

We need to know where data are stored when coding!

Pitfalls

- A **word** has four bytes
 - LW loads four bytes
 - There are four bytes in a word! They are located at sequential addresses
 - Sequential word addresses are incremented by 4!
- Sequential half words/bytes are **NOT** incremented by 4
 - Pay attention to the size
 - Sequential bytes do have sequential addresses
- Offset is a 12-bit 2's complement number, sign extended to 32 bits
 - If offset is too large, add offset with instructions
- Byte order matters

Summary of memory

- Memory is byte addressed
 - Each address identifies an 8-bit byte
 - A 32-bit address space support 4 GiB memory
- RV32I supports byte (8 bits), half-word (16 bits), and word (32 bits)
- Words and half-words should be aligned in memory
 - They must be aligned in this course
 - Although they do not have to in real processors, misalignment leads to poor performance
- Endianness affects the order of bytes when data are converted from/to bytes
 - RISC-V is little endian

Further thinking and reading

- How do you find out the endianness of a processor?
- Byte order is very important
 - Unicode BOM (byte order mark), U+FEFF
 - Search the Internet and find out how the mark is represented in UTF-16 (BE), UTF-16(LE), UTF-32(BE), and UTF-32(LE)

Find out what load/store instructions do

- Ask the following questions for load instructions
 - What is the address?
 - What are the bytes/is the byte the memory module finds at the address?
 - If there are multiple bytes, how should you put them together?
 - If necessary, how do you extend the byte(s) to 32 bits?
- Ask the following questions for store instructions
 - What is the address?
 - How many bytes are going to be stored in the address?
 - What is the order of bytes in the memory?

Question

What are the bits in t0 after the following instructions?

lh t0, 0x200(x0)

- A. 0x0000 81AC
- B. 0x0000 AC81
- C. 0xFFFF 81AC
- D. 0xFFFF AC81
- E. None of the above

Memory Address	Value
0x0000 0203	0x32
0x0000 0202	0x65
0x0000 0201	0x81
0x0000 0200	0xAC