

RISC-V Instruction Encoding - 1

32-bit Addr
MEM

0b 1001
↓
(5)₁₀

Data { 32
16
8
sign extension



UCONN

32-bit Instruction

Caiwen Ding

Department of Computer Science and Engineering
University of Connecticut

X () D

1000...0001...1111
Instruction

CSE3666: Introduction to Computer Architecture

area
↑
(860) 480-0006

admin

- Week 4
- Lab 3 is due on Wed.
- In this week's lectures, we will learn instruction encoding and function.
- HW2 is due on Friday.
- **Note:** In case you can not join the in-person class, we will use (<https://classrooms.uconn.edu/classroom/bous-a106/>) (BOUS A106 Remote access) provided by UConn office of the registrar to live stream the lectures.
- Without a **valid university accepted reason and evidence**, we will NOT be able to provide any further extension.
 - If you are unsure what valid university accepted reason and evidence is, please contact the **Deans of Students Office (DSO)** with your hardship in advance. DSO will send us an email for the record.
- When emailing TAs/Instructor, please add "3666" in the subject line such that we could better track your emails. We have 86 students enrolled; thus the response is inevitably slow. If you do not receive a reply in 48 hours, please email us again.

Outline

- RISC-V Instruction encoding
 - Instruction format → *several*
 - Encoding of instructions like ADD, ADDI, and Load/store
 - Decoding

Design Principle 4: Good design demands good compromises

Keep formats as similar as possible



Reading: Section 2.5 and the beginning of Section 2.10.

References: Reference card in the book.

Representing instructions with bits

- We use bits to represent numbers, characters, etc. ^{22-bit}
- We also use bits to represent instructions

Design questions:

00110011 ---- 1001 --- 0111 → machine code
?? ↗ Add t0, t1, t2

- How many bits should we use to encode instructions?
- Are we using the same number of bits to encode all instructions?
 - Do all instructions have the same length?

RISC-V instruction words

- RISC-V base ISA are encoded as **32-bit instruction words** *4 Bytes → 1 word*
 - Encoded instructions are also called machine (language) code
- Both instructions and data *Big* are stored in memory *Big* *??*
- ~~***~~ **Program Counter (PC)** points to the current instruction
 - ~~***~~ Incremented by 4 in normal flow for **sequential execution**

Do we have 1 Byte instruction? X
always 32-bit → 4 Bytes

Memory Address	Instructions
$x + 12$	Instr 13
$x + 8$	Instr 12
$x + 4$	Instr 11
x	Instr 10
$x - 4$...

virtually
 PC →

Bkpt	Address	Code	Basic	
<input type="checkbox"/>	0x00400000	0x0fc10513	laipc x10,0x0000fc10	16: la a0, msg
<input type="checkbox"/>	0x00400004	0x00050513	addi x10,x10,0x00000000	
<input type="checkbox"/>	0x00400008	0x00400893	addi x17,x0,0x00000004	17: li a7, 4
<input type="checkbox"/>	0x0040000c	0x00000073	ecall	18: ecall
<input type="checkbox"/>	0x00400010	0x00631293	slli x5,x6,0x00000006	19: slli t0,t1,6
<input type="checkbox"/>	0x00400014	0x00538333	add x6,x7,x5	20: loop: add t1, t2, t0
<input type="checkbox"/>	0x00400018	0x008000ef	jal x1,0x00000004	21: jal ra, foo
<input type="checkbox"/>	0x0040001c	0xfe031ce3	bne x6,x0,0xffffffffc	22: bne t1, zero, loop
<input type="checkbox"/>	0x00400020	0x00008067	jalr x0,x1,0x00000000	23: foo: jr ra
<input type="checkbox"/>	0x00400024	0x00a00893	addi x17,x0,0x0000000a	25: exit: li a7, 10
<input type="checkbox"/>	0x00400028	0x00000073	ecall	26: ecall

C extension allows compressed instructions of 16 bits, but it is not a stand-alone ISA.

Machine language uses binary representation of instructions.

Instructions in machine language are called **machine code**.

Discussion

- What information do you want to keep in the instruction word?

$$\boxed{7} \text{ bits} + \boxed{7} \text{ bits} + \boxed{1} + \boxed{7} = 32 \text{ bit}$$

ADD

rd, rs1, rs2

ADDI

rd, rs1, imm8

LW

rd, offset(rs1)

SW

rs2, offset(rs1)

instruction

32 bits

Instruction format

- The layout of bits in instruction words is **instruction format**
 - How do we use 32-bit bits to specify operation code (opcode), registers, immediate, offset, etc? How many bits for each?
- RISC-V has six instruction formats (while MIPS has 3)
 - R, I, S, SB, U, and UJ
- Instructions we have learned so far
 - ? – Using registers only *add, sll*
 - ? – Having an immediate as the second operand *addi, slli*
 - ? – Load and store *lw, sw*
 - ? – LUI
 - ? – Branch, to be discussed next week

Binary compatibility allows compiled programs to work on different computers

RISC-V Core Instruction Format



	31	27	26	25	24	20	19	15	14	12	11	7	6	0
R	funct7					rs2		rs1		funct3		rd		Opcode
I	imm[11:0]					rs1		funct3		rd		Opcode		
S	imm[11:5]					rs2		rs1		funct3		imm[4:0]		opcode
SB	imm[12 10:5]					rs2		rs1		funct3		imm[4:1 11]		opcode
U	imm[31:12]										rd		opcode	
UJ	imm[20 10:1 11 19:12]										rd		opcode	

RISC-V Reference Data Card ("Green Card")

```

alt      R Set Less Than          R[rd] = (R[rs1] < R[rs2]) ? 1 : 0
altiu    I Set Less Than Immediate R[rd] = (R[rs1] < imm) ? 1 : 0
altiu    I Set < Immediate Unsigned R[rd] = (R[rs1] < imm) ? 1 : 0
altiu    R Set Less Than Unsigned  R[rd] = (R[rs1] < R[rs2]) ? 1 : 0
sra,sraw R Shift Right Arithmetic (Word) R[rd] = R[rs1] >> R[rs2]
sra,sraw I Shift Right Arithmetic (Word) R[rd] = R[rs1] >> imm
srli,srli R Shift Right (Word)          R[rd] = R[rs1] >> R[rs2]
srli,srli I Shift Right (Word)          R[rd] = R[rs1] >> imm
srlw,srlw R Shift Right Arithmetic (Word) R[rd] = R[rs1] >> R[rs2]
srlw,srlw I Shift Right Arithmetic (Word) R[rd] = R[rs1] >> imm
sub,subw R Subtract (Word)            R[rd] = R[rs1] - R[rs2]
sub,subw I Subtract (Word)            R[rd] = R[rs1] - R[rs2]
sw       S Store Word                 M[rs1 + imm(31:0)] = R[rs2(31:0)]
xor      R XOR                       R[rd] = R[rs1] ^ R[rs2]
xor      I XOR Immediate             R[rd] = R[rs1] ^ imm
  
```

Notes: 1) The Word version only operates on the rightmost 32 bits of a 64-bit registers
 2) Operation assumes unsigned integers (instead of 2's complement)
 3) The least significant bit of the branch address in jalr is set to 0
 4) (signed) Load instructions extend the sign bit of data to fill the 64-bit register
 5) Replicates the sign bit to fill in the leftmost bits of the result during right shift
 6) Multiply with one operand signed and one unsigned
 7) The Single version does a single-precision operation using the rightmost 32 bits of a 64-bit F register
 8) Classify writes a 10-bit mask to show which properties are true (e.g., -inf, -0, +0, +inf, denorm, ...)
 9) Atomic memory operation; nothing else can interpose itself between the read and the write of the memory location
 The immediate field is sign-extended in RISC-V

```

amoxxor.w,amoxxor.d R MAXimum Unsigned  if (R[rs2] > M[R[rs1]]) M[R[rs1]] = R[rs2]
amoxxor.w,amoxxor.d R MINimum           if (R[rs2] > M[R[rs1]]) M[R[rs1]] = R[rs2]
amoxxor.w,amoxxor.d R MINimum Unsigned  if (R[rs2] < M[R[rs1]]) M[R[rs1]] = R[rs2]
amoxxor.w,amoxxor.d R OR                 R[rd] = M[R[rs1]]
amoxxor.w,amoxxor.d R OR                 R[rd] = M[R[rs1]]
amoxxor.w,amoxxor.d R SWAP               M[R[rs1]] = M[R[rs1]] | R[rs2]
amoxxor.w,amoxxor.d R XOR                R[rd] = M[R[rs1]] ^ M[R[rs1]] ^ R[rs2]
amoxxor.w,amoxxor.d R XOR                R[rd] = M[R[rs1]] ^ R[rs2]
lr.w,lr.d           R Load Reserved     reservation on M[R[rs1]]
sc.w,sc.d           R Store Conditional if reserved, M[R[rs1]] = R[rs2],
                                     R[rd] = 0, else R[rd] = 1
  
```


CORE INSTRUCTION FORMATS

	31	27	26	25	24	20	19	15	14	12	11	7	6	0						
R	funct7					rs2		rs1		funct3		rd		Opcode						
I	imm[11:0]					rs1		funct3		rd		Opcode								
S	imm[11:5]					rs2		rs1		funct3		imm[4:0]		opcode						
SB	imm[12 10:5]					rs2		rs1		funct3		imm[4:1 11]		opcode						
U	imm[31:12]											rd		opcode						
UJ	imm[20 10:1 11 19:12]											rd		opcode						

The (green) card in the textbook

imm is the 32-bit immediate processor uses for computation.
 It is not the immediate written in instruction.

RISC-V R-type Instructions



- For instructions that have 3 registers as operands
- Fields in R-type
 - opcode: 7-bit operation code
 - rd: destination register number
 - rs1: first source register number
 - rs2: second source register number
 - funct3: additional function code
 - funct7: even more function code

R-format Example: ADD

funct7	rs2	rs1	funct3	rd	opcode
--------	-----	-----	--------	----	--------

7 bits

5 bits

5 bits

3 bits

5 bits

7 bits

rd rs1 rs2
add x1, x2, x3

<i>0000000</i>	<i>00011</i>	<i>00010</i>	<i>000</i>	<i>00001</i>	<i>0110011</i>
----------------	--------------	--------------	------------	--------------	----------------

x3
△

x2
△

x1
△

Opcode and funct codes

- From the green card, which also has hexadecimal representation

	type	opcode	funct3	funct7
add	R	0110011	000	0000000
sub	R	0110011	000	0100000
sll	R	0110011	001	0000000
slt	R	0110011	010	0000000
sltu	R	0110011	011	0000000
xor	R	0110011	100	0000000
srl	R	0110011	101	0000000
sra	R	0110011	101	0100000
or	R	0110011	110	0000000
and	R	0110011	111	0000000

The right most two bits in opcode are always 11
for 32-bit instructions

R-type Example: ADD

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

~~add~~ x1, x2, x3

Sub

0	3	2	0	1	0x33
---	---	---	---	---	------

000 0000	00011	00010	000	00001	011 0011
----------	-------	-------	-----	-------	----------

Sub ?

0x 003100B3

What if we change ADD to SUB? How many bits are changed?

Question (from textbook)


- What RISC-V instruction does this represent? Choose from one of the four options below?
 - The numbers are in decimal

funct7	rs2	rs1	funct3	rd	opcode
32	9	10	0	11	51

- A. sub x9, x10, x11
- B. add x11, x9, x10
- C. sub x11, x10, x9
- D. sub x11, x9, x10

	Funct7
ADD	0b 000 0000
SUB	0b 010 0000

RISC-V R-type Instructions



imm
5 bits ??

- We can use the format to encode any instructions like

instr_name rd, rs1, rs2

- How about instructions like addi, slli?

addi rd, rs1, imm
12 bits
range

RISC-V I-type Instructions



- Fields in I-type
 - opcode: operation code
 - rd: destination register number
 - rs1: first source register number
 - funct3: additional function code
 - imm: lower 12 bits of the immediate, in the place of funct7 and rs2
- Since only 12 bits are kept in machine code, the immediate must be in $[-2^{11}, +2^{11} - 1]$ or $[-2048, 2047]$

Example of I-type opcode and funct3 code


	type	opcode	funct3
<i>in</i> <u>addi</u>	I	<u>0010011</u>	000
<u>slli</u>	I	0010011	001
slti	I	0010011	010
sltiu	I	0010011	011
xori	I	0010011	100
<u>srli</u>	I	0010011	<u>101</u>
srai	I	0010011	<u>101</u>
ori	I	<u>0010011</u>	<u>110</u>
andi	I	0010011	<u>111</u>

In **slli**, **srli**, and **srai**, only lower 5 bits of the immediate are used for shift amount.
5 bits are enough for 32-bit registers!

I-type Example: ADDI

imm[11:0]	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits

rd rs1
addi x1, x2, 32

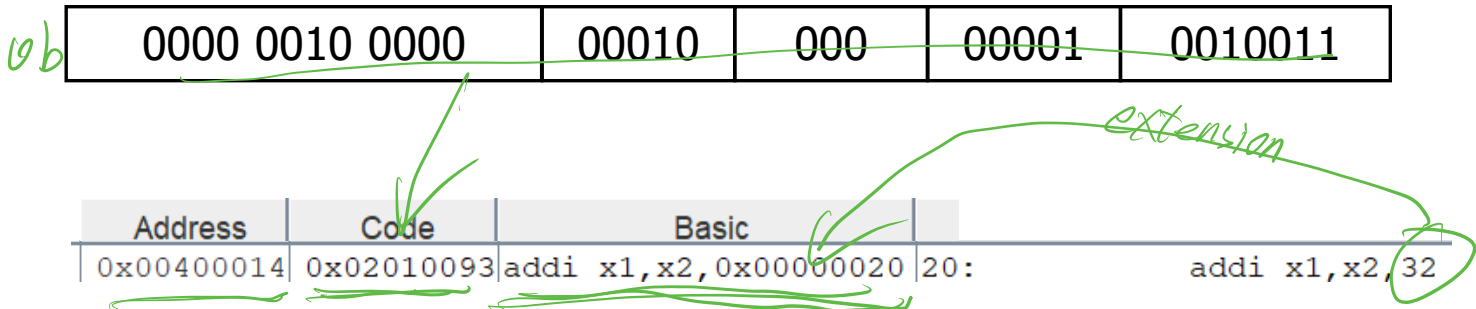


0 0000 0000	00010	000	00001	0010011
	<i>x2</i>		<i>rd</i>	

I-type Example: ADDI

imm[11:0]	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits

`addi x1, x2, 32`



When executing the instruction, processor builds a 32-bit immediate **imm**



I-type Example: SRLI vs SRAI

imm[11:0]	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits

srli x1, x2, 16



0000 0001 0000	00010	101	00001	0010011
----------------	-------	-----	-------	---------



2nd MSB

srai x1, x2, 16



Same opcode and funct3
Bit 30 is different!

0100 0001 0000	00010	101	00001	0010011
----------------	-------	-----	-------	---------



srli	I	0010011	101	0000000	← funct7
srai	I	0010011	101	0100000	

What format would you use for load instructions?

- Load instructions

lw rd, offset(rs1)


imm

A. R-format

B. I-format

R:	funct7	rs2	rs1	funct3	rd	opcode
I:	imm [11:0]		rs1	funct3	rd	opcode

Loads are I-type



lw x1, 32(x2)



sw x2, 4(x1) ? → rd??

0000 0010 0000	00010	010	00001	000 0011
----------------	-------	-----	-------	----------

opcode

funct3

lb	I	0000011	000
lh	I	0000011	001
lw	I	0000011	010
ld	I	0000011	011
lbu	I	0000011	100
lhu	I	0000011	101
lwu	I	0000011	110

Do you see the pattern in funct3?

How about store?


How about store instructions?

- Store instructions have rs2, but not rd

`sw rs2,offset(rs1)`

R:	funct7	rs2	rs1	funct3	rd	opcode
I:	imm [11:0]		rs1	funct3	rd	opcode

RISC-V S-type Instructions



- Fields in S-type


- opcode: operation code
- rs1: first source register number
- rs2: second source register number
- imm[11:5] and imm[4:0]:

The lower 12 bits of the offset/immediate are stored into two fields.

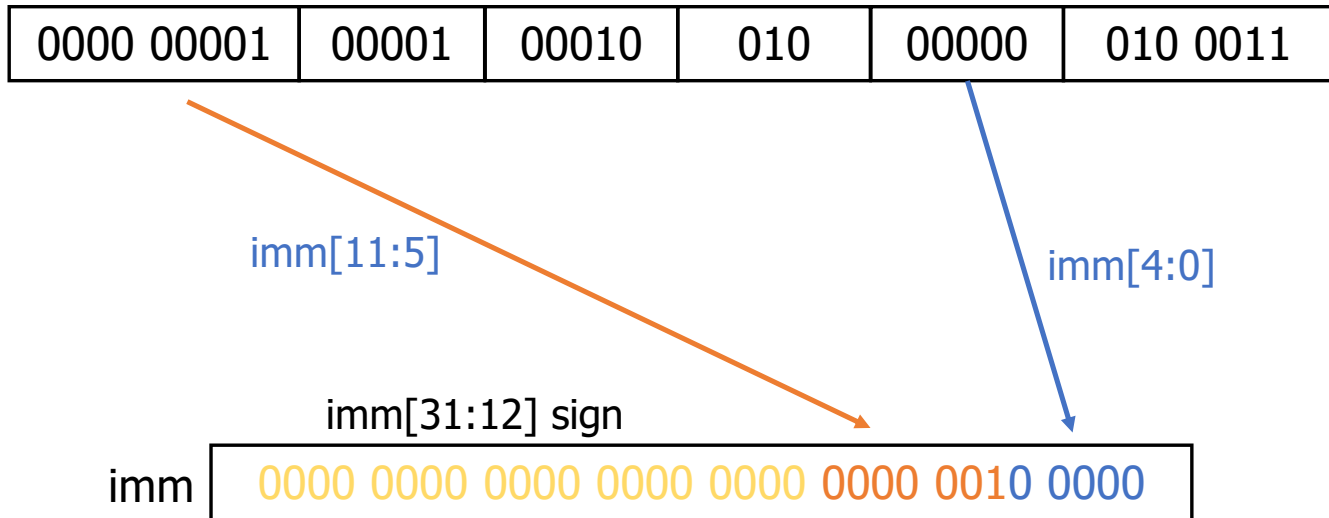
The higher 7 bits, bits 11 to 5, are in funct7

The lower 5 bits, bits 4 to 0, are in rd

Stores are S-type



sw x1, 32(x2) The lower 12 bits of the immediate are saved in two fields



Summary of R-, I-, and S-type instructions

Instruction	Format	funct7	rs2	rs1	funct3	rd	opcode
add (add)	R	0000000	reg	reg	000	reg	0110011
sub (sub)	R	0100000	reg	reg	000	reg	0110011
Instruction	Format	immediate		rs1	funct3	rd	opcode
addi (add immediate)	I	constant		reg	000	reg	0010011
lw (load word)	I	address		reg	010	reg	0000011
Instruction	Format	immediate	rs2	rs1	funct3	immediate	opcode
sw (store word)	S	address	reg	reg	010	address	0100011

Fields, other than immediate fields, are located at the same location, for all types
 Placement of bits in the immediate is more complicated

Examples of R-, I-, and S-type instructions

R-type Instructions	funct7	rs2	rs1	funct3	rd	opcode	Example
add (add)	0000000	00011	00010	000	00001	0110011	add x1, x2, x3
sub (sub)	0100000	00011	00010	000	00001	0110011	sub x1, x2, x3
I-type Instructions	immediate		rs1	funct3	rd	opcode	Example
addi (add immediate)	001111101000		00010	000	00001	0010011	addi x1, x2, 1000
lw (load word)	001111101000		00010	010	00001	0000011	lw x1, 1000 (x2)
S-type Instructions	immed-iate	rs2	rs1	funct3	immed-iate	opcode	Example
sw (store word)	0011111	00001	00010	010	01000	0100011	sw x1, 1000(x2)

RISC-V Core Instruction Format

- Now we know four types: R, I, S, and U
- We will discuss SB, and UJ later

	31	27	26	25	24	20	19	15	14	12	11	7	6	0
R	funct7				rs2		rs1		funct3		rd		Opcode	
I	imm[11:0]						rs1		funct3		rd		Opcode	
S	imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode	
SB	imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode	
U	imm[31:12]										rd		opcode	
UJ	imm[20 10:1 11 19:12]										rd		opcode	

Only real instructions can be encoded into machine code.

Exercise

- Pick an instruction and encode it
- Study the machine code generated by RARS

Example

Hex: 00030503

Bin: 0000 0000 0000 0011 0000 0101 0000 0011

Fields: 0000 0000 0000 0011 0000 0101 0000 0011

lb	I	0000011	000
lh	I	0000011	001
lw	I	0000011	010
ld	I	0000011	011
lbu	I	0000011	100
lhu	I	0000011	101
lwu	I	0000011	110

Questions

- RISC-V has three fields for specifying operations: opcode, funct3, and funct7. Why don't they combine them and have a single opcode field of 17 bits?
- If someone decides to increase the number of registers to 64, how does it affect the encoding of instructions?